

Theoretical

```
In [ ]: #1. What is Boosting in Machine Learning?
'''
Boosting is an ensemble learning technique that improves prediction accuracy by com
Each model corrects the errors of the previous one by focusing more on the misclass
the final model more robust. Boosting is widely used in both classification and reg
performance on complex datasets. Some popular Boosting algorithms include AdaBoost,
'''

#2. How does Boosting differ from Bagging?
'''
Bagging and Boosting are both ensemble techniques, but they work differently. Baggi
parallel using bootstrapped datasets and averages their predictions to reduce varia
models sequentially, where each model learns from the mistakes of the previous one,
reducing overfitting, while Boosting is useful for improving accuracy on complex da
algorithm, whereas AdaBoost and Gradient Boosting are examples of Boosting.
'''

#3. What is the key idea behind AdaBoost?
'''
AdaBoost (Adaptive Boosting) is a Boosting algorithm that combines multiple weak le
assigns higher weights to misclassified samples, making subsequent models focus mor
iteration, a new weak learner is trained on the updated weights, and their predicti
AdaBoost is mainly used with decision stumps (one-level decision trees) but can be
widely used in face detection, fraud detection, and text classification.
'''

#4. Explain the working of AdaBoost with an example.
'''
AdaBoost starts by assigning equal weights to all training samples and training a w
If a sample is misclassified, its weight is increased so that the next weak classif
is repeated for several iterations, with each new classifier improving upon the err
classifiers' predictions are combined using weighted voting. For example, in a spam
iteratively refine its decision rules, focusing on emails that were previously misc
'''

#5. What is Gradient Boosting, and how is it different from AdaBoost?
'''
Gradient Boosting is an advanced Boosting technique that builds models sequentially
Instead of adjusting sample weights, it minimizes the residual errors using gradien
predict the residuals (errors) of the previous model, gradually improving accuracy.
to training samples, Gradient Boosting directly optimizes the loss function. This m
effective, especially for regression tasks and high-dimensional datasets.
'''

#6. What is the Loss function in Gradient Boosting?
'''
Gradient Boosting minimizes a specified loss function to improve predictions. For r
include Mean Squared Error (MSE) and Mean Absolute Error (MAE), which measure diffe
For classification, Log Loss (Cross-Entropy) is commonly used, ensuring that the pr
loss function guides how new weak learners are added to the model. By minimizing th
Boosting effectively reduces errors and enhances model accuracy.
'''
```

#7. How does XGBoost improve over traditional Gradient Boosting?

'''

XGBoost (Extreme Gradient Boosting) enhances traditional Gradient Boosting by introducing parallel processing to speed up training and implements L1/L2 regularization to prevent overfitting. It also features a more efficient tree-pruning strategy, handling missing values automatically and allowing for better feature selection. XGBoost is highly scalable, making it a preferred choice for large datasets. It is widely used in applications like recommendation systems, finance, and more.

'''

#8. What is the difference between XGBoost and CatBoost?

'''

XGBoost and CatBoost are both Gradient Boosting-based algorithms, but they handle categorical variables differently. XGBoost requires categorical variables to be manually encoded using one-hot encoding or label encoding. CatBoost, on the other hand, is specifically designed for categorical data and uses Ordered Target Encoding, which is more efficient. CatBoost also incorporates a unique ordered boosting technique, which helps in reducing overfitting. CatBoost is often preferred for datasets with many categorical features, such as customer data.

'''

#9. What are some real-world applications of Boosting techniques?

'''

Boosting techniques are widely used across various domains due to their high accuracy and ability to handle complex, non-linear relationships. Some real-world applications include credit scoring and fraud detection in finance, disease prediction in healthcare, product recommendations in e-commerce, and spam detection in email services. Boosting models are also popular in autonomous driving for object detection and classification.

'''

#10. How does regularization help in XGBoost?

'''

Regularization in XGBoost helps in controlling overfitting by penalizing complex models. It includes L1 (Lasso) and L2 (Ridge) regularization to shrink feature weights and prevent overly complex trees. This ensures that the model generalizes better to new data. Additionally, XGBoost includes parameters like `max_depth` and `min_child_weight` to further control model complexity and prevent overfitting. These regularization techniques make XGBoost more stable and suitable for large, noisy datasets.

'''

#11. What are some hyperparameters to tune in Gradient Boosting models?

'''

Tuning hyperparameters in Gradient Boosting is essential for optimal performance. Key parameters include `learning_rate` (controls the contribution of each weak learner), `n_estimators` (number of estimators), `max_depth` (maximum depth of trees to prevent overfitting), `subsample` (controls the fraction of samples used for each tree), and `colsample_bytree` (controls the fraction of features used for each tree). Proper tuning ensures the best trade-off between bias and variance.

'''

#12. What is the concept of Feature Importance in Boosting?

'''

Feature Importance in Boosting helps identify the most influential features in a dataset. It assigns scores to features based on their contribution to improving predictions. Features with higher scores have a greater impact on the model's performance. Feature Importance is useful for feature selection, identifying redundant features to be removed for better model efficiency, and interpreting model decisions.

'''

#13. Why is CatBoost efficient for categorical data?

'''

CatBoost is designed to handle categorical variables efficiently without requiring Target Encoding, which prevents data leakage by encoding categories based on past data. This technique reduces overfitting by creating more generalized models. Unlike XGBoost, CatBoost automatically processes categorical features, making it ideal for datasets like customer segmentation and sentiment analysis.

Practical

```
In [1]: #Train an AdaBoost Classifier on a sample dataset and print model accuracy
from sklearn.ensemble import AdaBoostClassifier
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.tree import DecisionTreeClassifier

X, y = make_classification(n_samples=1000, n_features=20, random_state=42)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

base_estimator = DecisionTreeClassifier(max_depth=1)
adaboost_clf = AdaBoostClassifier(base_estimator, n_estimators=50, random_state=42)

adaboost_clf.fit(X_train, y_train)

y_pred = adaboost_clf.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
print(f"Model Accuracy: {accuracy:.2f}")
```

Model Accuracy: 0.87

```
In [3]: #Train an AdaBoost Regressor and evaluate performance using Mean Absolute Error (MAE)
from sklearn.datasets import make_regression
from sklearn.ensemble import AdaBoostRegressor
from sklearn.metrics import mean_absolute_error
from sklearn.tree import DecisionTreeRegressor

X, y = make_regression(n_samples=1000, n_features=20, noise=0.1, random_state=42)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

base_estimator = DecisionTreeRegressor(max_depth=4)
adaboost_reg = AdaBoostRegressor(base_estimator, n_estimators=50, random_state=42)

adaboost_reg.fit(X_train, y_train)
```

```

y_pred = adaboost_reg.predict(X_test)

mae = mean_absolute_error(y_test, y_pred)
print(f"Mean Absolute Error (MAE): {mae:.2f}")

```

Mean Absolute Error (MAE): 73.09

```

In [5]: #Train a Gradient Boosting Classifier on the Breast Cancer dataset and print feature importance
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.datasets import load_breast_cancer
import pandas as pd

data = load_breast_cancer()
X, y = data.data, data.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

gb_clf = GradientBoostingClassifier(n_estimators=100, learning_rate=0.1, random_state=42)
gb_clf.fit(X_train, y_train)

feature_importance = gb_clf.feature_importances_

feature_importance_df = pd.DataFrame(
    {"Feature": data.feature_names, "Importance": feature_importance}
).sort_values(by="Importance", ascending=False)

print(feature_importance_df)

```

	Feature	Importance
7	mean concave points	0.450418
27	worst concave points	0.240209
20	worst radius	0.075424
22	worst perimeter	0.051441
21	worst texture	0.039881
23	worst area	0.038200
1	mean texture	0.027821
26	worst concavity	0.017576
16	concavity error	0.012933
13	area error	0.010848
10	radius error	0.005238
24	worst smoothness	0.004811
19	fractal dimension error	0.004313
5	mean compactness	0.003838
11	texture error	0.003488
15	compactness error	0.002644
17	concave points error	0.002072
4	mean smoothness	0.002039
28	worst symmetry	0.001500
6	mean concavity	0.001115
25	worst compactness	0.000891
18	symmetry error	0.000709
14	smoothness error	0.000669
3	mean area	0.000555
8	mean symmetry	0.000466
12	perimeter error	0.000356
9	mean fractal dimension	0.000317
2	mean perimeter	0.000201
0	mean radius	0.000013
29	worst fractal dimension	0.000013

```
In [6]: #4 Train a Gradient Boosting Regressor and evaluate using R-Squared Score
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.metrics import r2_score

X, y = make_regression(n_samples=1000, n_features=20, noise=0.1, random_state=42)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

gb_reg = GradientBoostingRegressor(n_estimators=100, learning_rate=0.1, random_state=42)
gb_reg.fit(X_train, y_train)

y_pred = gb_reg.predict(X_test)

r2 = r2_score(y_test, y_pred)
print(f"R-Squared Score: {r2:.2f}")
```

R-Squared Score: 0.92

```
In [ ]: !pip install xgboost
```

Defaulting to user installation because normal site-packages is not writeable
 Looking in links: /usr/share/pip-wheels
 Collecting xgboost
 Downloading xgboost-3.0.0-py3-none-manylinux_2_28_x86_64.whl.metadata (2.1 kB)
 Requirement already satisfied: numpy in /opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-packages (from xgboost) (1.26.4)
 Collecting nvidia-nccl-cu12 (from xgboost)
 Downloading nvidia_nccl_cu12-2.26.2-py3-none-manylinux2014_x86_64.manylinux_2_17_x86_64.whl.metadata (2.0 kB)
 Requirement already satisfied: scipy in /opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-packages (from xgboost) (1.12.0)
 Downloading xgboost-3.0.0-py3-none-manylinux_2_28_x86_64.whl (253.9 MB)
 253.9/253.9 MB 28.5 MB/s eta 0:00:00:0
 0:01[36m0:00:01
 Downloading nvidia_nccl_cu12-2.26.2-py3-none-manylinux2014_x86_64.manylinux_2_17_x86_64.whl (201.3 MB)
 201.3/201.3 MB 80.5 MB/s eta 0:00:00:0
 0:01[36m0:00:01
 Installing collected packages: nvidia-nccl-cu12, xgboost

```
In [10]: #Train an XGBoost Classifier on a dataset and compare accuracy with Gradient Boosting
from sklearn.ensemble import GradientBoostingClassifier
from xgboost import XGBClassifier
from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score

data = load_breast_cancer()
X, y = data.data, data.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

gb_clf = GradientBoostingClassifier(n_estimators=100, learning_rate=0.1, random_state=42)
gb_clf.fit(X_train, y_train)

xgb_clf = XGBClassifier(n_estimators=100, learning_rate=0.1, eval_metric="logloss", random_state=42)
xgb_clf.fit(X_train, y_train)

y_pred_gb = gb_clf.predict(X_test)
y_pred_xgb = xgb_clf.predict(X_test)

accuracy_gb = accuracy_score(y_test, y_pred_gb)
accuracy_xgb = accuracy_score(y_test, y_pred_xgb)

print(f"Gradient Boosting Accuracy: {accuracy_gb:.4f}")
print(f"XGBoost Accuracy: {accuracy_xgb:.4f}")
```

Gradient Boosting Accuracy: 0.9561
 XGBoost Accuracy: 0.9561

In [11]: `!pip install catboost`

```

Defaulting to user installation because normal site-packages is not writeable
Looking in links: /usr/share/pip-wheels
Collecting catboost
  Downloading catboost-1.2.7-cp310-cp310-manylinux2014_x86_64.whl.metadata (1.2 kB)
Requirement already satisfied: graphviz in /home/104c0660-1be7-404b-8861-82ee06d5989e/.local/lib/python3.10/site-packages (from catboost) (0.20.3)
Requirement already satisfied: matplotlib in /opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-packages (from catboost) (3.8.0)
Requirement already satisfied: numpy<2.0,>=1.16.0 in /opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-packages (from catboost) (1.26.4)
Requirement already satisfied: pandas>=0.24 in /opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-packages (from catboost) (2.1.4)
Requirement already satisfied: scipy in /opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-packages (from catboost) (1.12.0)
Requirement already satisfied: plotly in /opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-packages (from catboost) (5.19.0)
Requirement already satisfied: six in /opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-packages (from catboost) (1.16.0)
Requirement already satisfied: python-dateutil>=2.8.2 in /opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-packages (from pandas>=0.24->catboost) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-packages (from pandas>=0.24->catboost) (2023.3.post1)
Requirement already satisfied: tzdata>=2022.1 in /opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-packages (from pandas>=0.24->catboost) (2023.3)
Requirement already satisfied: contourpy>=1.0.1 in /opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-packages (from matplotlib->catboost) (1.2.0)
Requirement already satisfied: cyclor>=0.10 in /opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-packages (from matplotlib->catboost) (0.11.0)
Requirement already satisfied: fonttools>=4.22.0 in /opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-packages (from matplotlib->catboost) (4.25.0)
Requirement already satisfied: kiwisolver>=1.0.1 in /opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-packages (from matplotlib->catboost) (1.4.4)
Requirement already satisfied: packaging>=20.0 in /opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-packages (from matplotlib->catboost) (23.2)
Requirement already satisfied: pillow>=6.2.0 in /opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-packages (from matplotlib->catboost) (10.2.0)
Requirement already satisfied: pyparsing>=2.3.1 in /opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-packages (from matplotlib->catboost) (3.0.9)
Requirement already satisfied: tenacity>=6.2.0 in /opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-packages (from plotly->catboost) (8.2.2)
Downloading catboost-1.2.7-cp310-cp310-manylinux2014_x86_64.whl (98.7 MB)
----- 98.7/98.7 MB 58.9 MB/s eta 0:00:00 0:00:01
01[36m0:00:01
Installing collected packages: catboost
Successfully installed catboost-1.2.7

```

In [12]: `#Train a CatBoost Classifier and evaluate using F1-Score`

```

from catboost import CatBoostClassifier
from sklearn.datasets import load_breast_cancer
from sklearn.metrics import f1_score

```

```

data = load_breast_cancer()
X, y = data.data, data.target

```

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

cat_clf = CatBoostClassifier(n_estimators=100, learning_rate=0.1, verbose=0, random_state=42)
cat_clf.fit(X_train, y_train)

y_pred_cat = cat_clf.predict(X_test)

f1 = f1_score(y_test, y_pred_cat)

print(f"F1-Score: {f1:.4f}")

```

F1-Score: 0.9722

```

In [13]: #Train an XGBoost Regressor and evaluate using Mean Squared Error (MSE)
from xgboost import XGBRegressor
from sklearn.datasets import make_regression
from sklearn.metrics import mean_squared_error

X, y = make_regression(n_samples=1000, n_features=20, noise=0.1, random_state=42)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

xgb_reg = XGBRegressor(n_estimators=100, learning_rate=0.1, random_state=42)
xgb_reg.fit(X_train, y_train)

y_pred_xgb = xgb_reg.predict(X_test)

mse = mean_squared_error(y_test, y_pred_xgb)

print(f"Mean Squared Error: {mse:.4f}")

```

Mean Squared Error: 4845.6934

```

In [15]: #Train an AdaBoost Classifier and visualize feature importance
import matplotlib.pyplot as plt

data = load_breast_cancer()
X, y = data.data, data.target
feature_names = data.feature_names

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

ada_clf = AdaBoostClassifier(n_estimators=100, learning_rate=0.1, random_state=42)

```



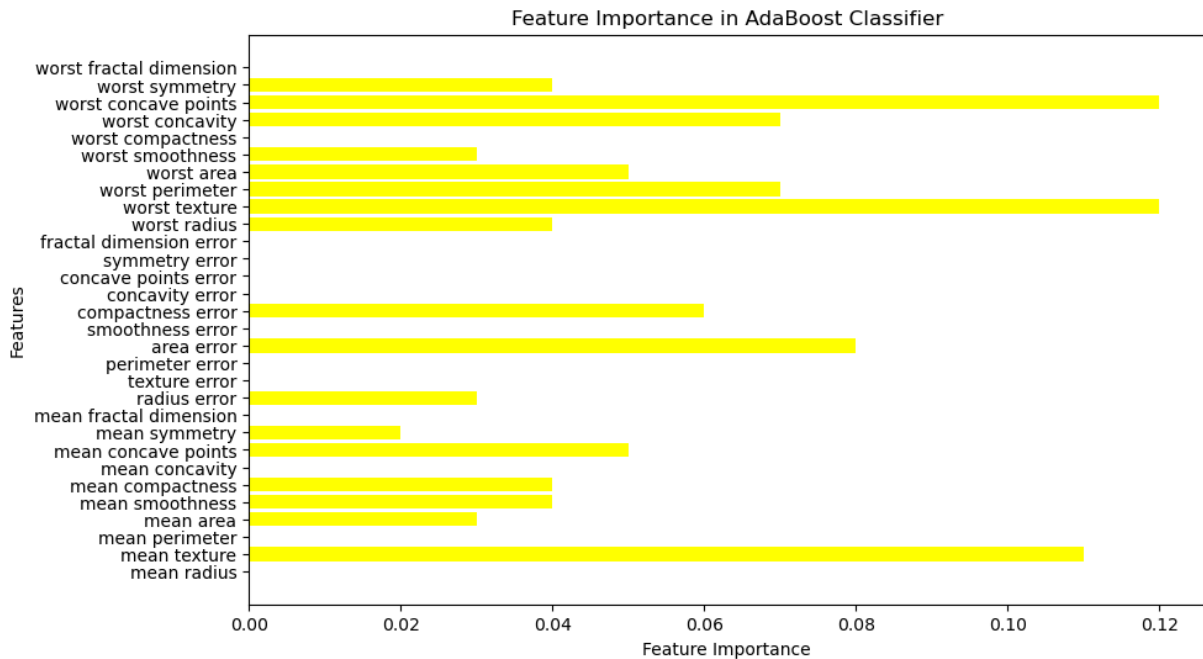
```

ada_clf.fit(X_train, y_train)

feature_importance = ada_clf.feature_importances_

plt.figure(figsize=(10, 6))
plt.barh(feature_names, feature_importance, color='yellow')
plt.xlabel("Feature Importance")
plt.ylabel("Features")
plt.title("Feature Importance in AdaBoost Classifier")
plt.show()

```



```

In [17]: #Train a Gradient Boosting Regressor and plot Learning curves

import numpy as np

X, y = make_regression(n_samples=1000, n_features=20, noise=0.1, random_state=42)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize the Gradient Boosting Regressor
gb_reg = GradientBoostingRegressor(n_estimators=100, learning_rate=0.1, random_state=42)

train_errors = []
test_errors = []

for n_estimators in range(1, 101):
    gb_reg.set_params(n_estimators=n_estimators)
    gb_reg.fit(X_train, y_train)

    y_train_pred = gb_reg.predict(X_train)

```

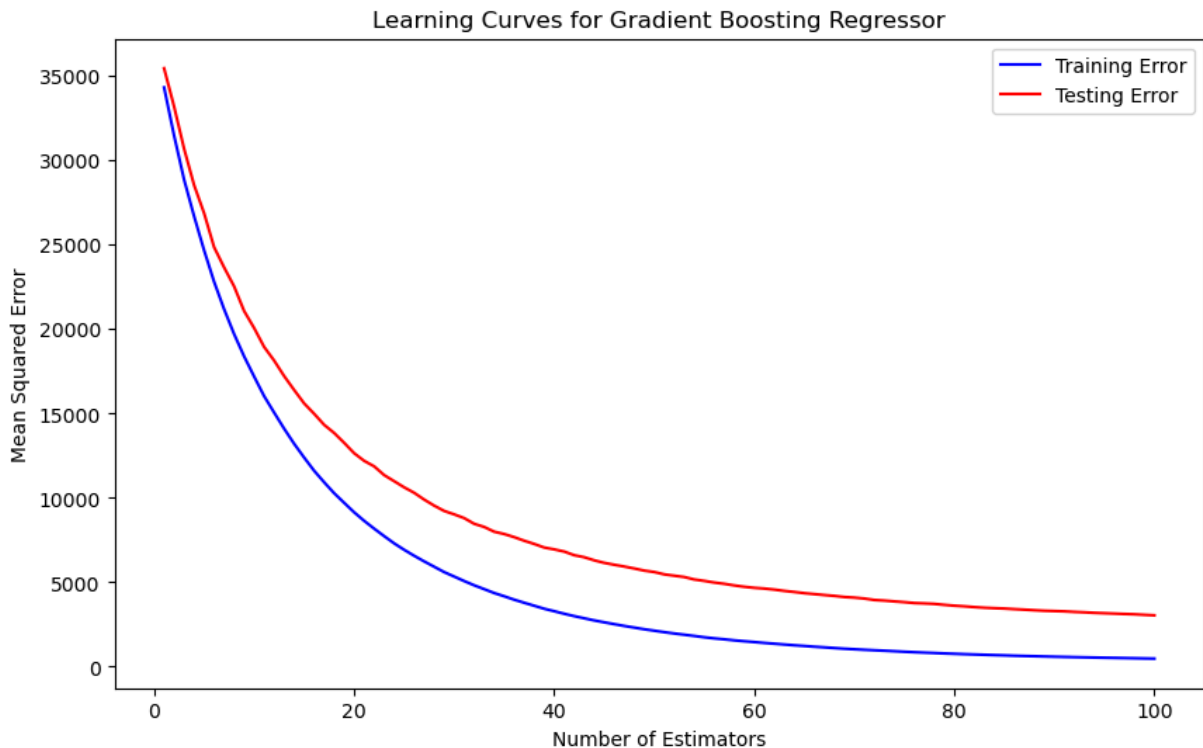
```

y_test_pred = gb_reg.predict(X_test)

train_errors.append(mean_squared_error(y_train, y_train_pred))
test_errors.append(mean_squared_error(y_test, y_test_pred))

plt.figure(figsize=(10, 6))
plt.plot(range(1, 101), train_errors, label="Training Error", color="blue")
plt.plot(range(1, 101), test_errors, label="Testing Error", color="red")
plt.xlabel("Number of Estimators")
plt.ylabel("Mean Squared Error")
plt.title("Learning Curves for Gradient Boosting Regressor")
plt.legend()
plt.show()

```



In [18]: *#Train an XGBoost Classifier and visualize feature importances*

```

data = load_breast_cancer()
X, y = data.data, data.target
feature_names = data.feature_names

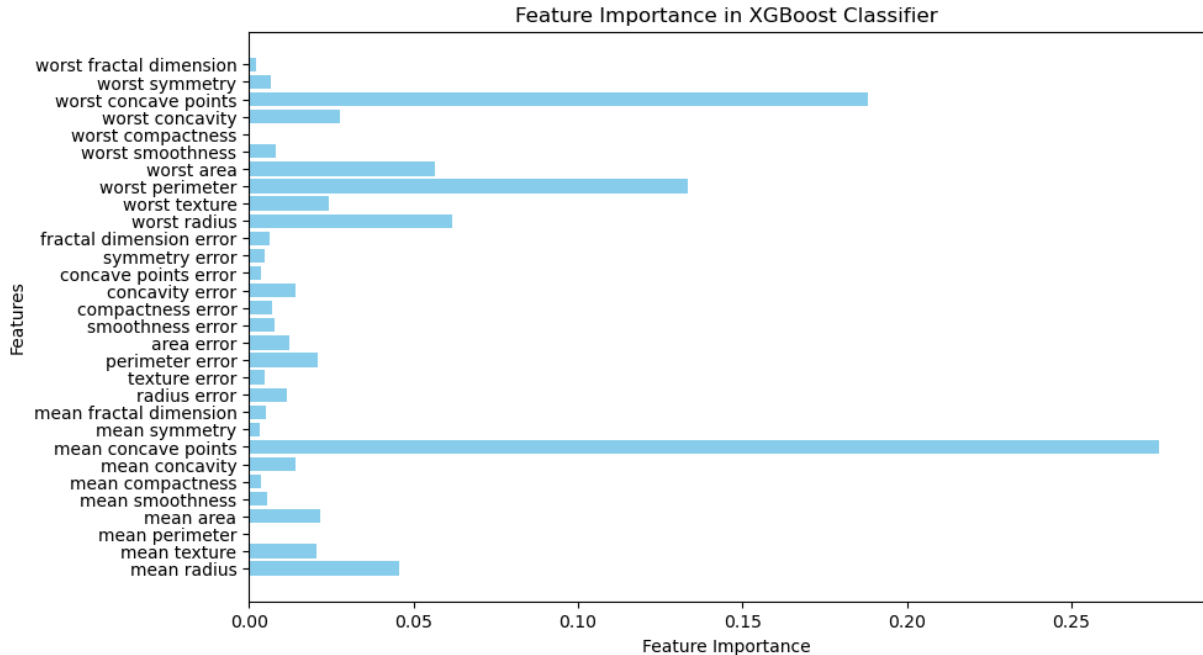
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

xgb_clf = XGBClassifier(n_estimators=100, learning_rate=0.1, random_state=42)
xgb_clf.fit(X_train, y_train)

feature_importance = xgb_clf.feature_importances_

```

```
plt.figure(figsize=(10, 6))
plt.barh(feature_names, feature_importance, color='skyblue')
plt.xlabel("Feature Importance")
plt.ylabel("Features")
plt.title("Feature Importance in XGBoost Classifier")
plt.show()
```



```
In [19]: #Train a CatBoost Classifier and plot the confusion matrix
import seaborn as sns
from sklearn.metrics import confusion_matrix

data = load_breast_cancer()
X, y = data.data, data.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

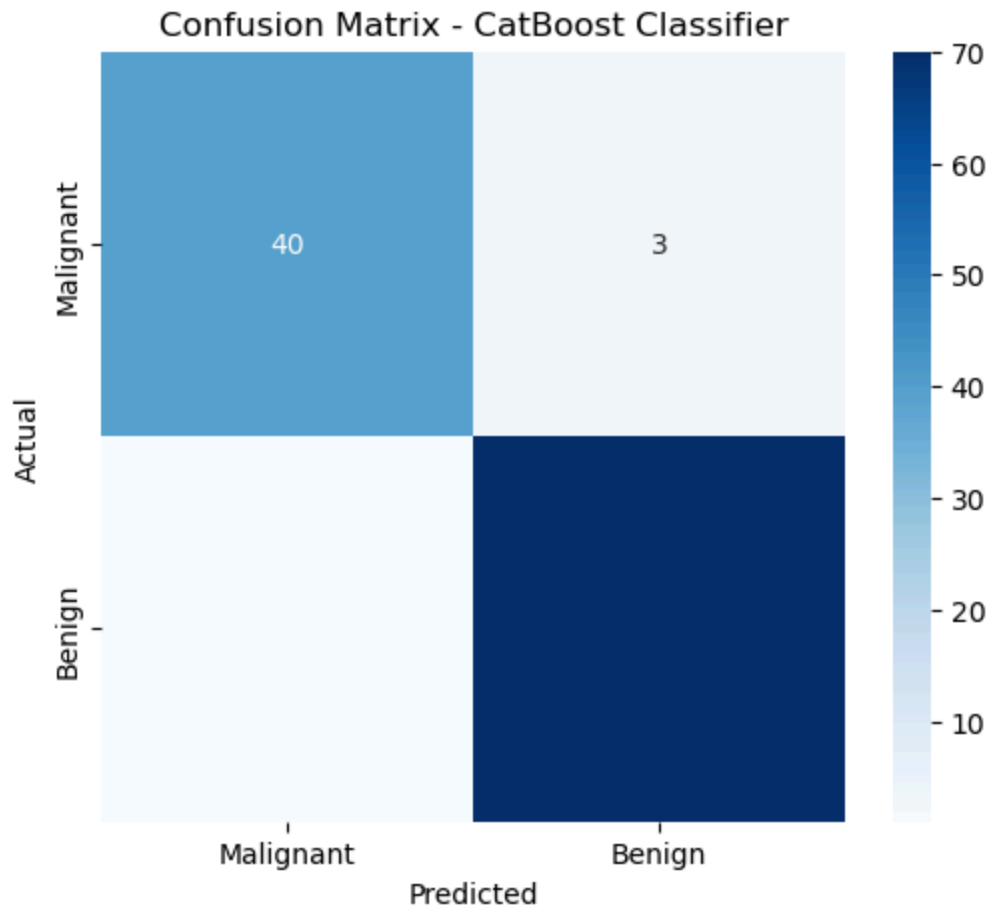
cat_clf = CatBoostClassifier(n_estimators=100, learning_rate=0.1, verbose=0, random_state=42)
cat_clf.fit(X_train, y_train)

y_pred_cat = cat_clf.predict(X_test)

cm = confusion_matrix(y_test, y_pred_cat)

plt.figure(figsize=(6, 5))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=["Malignant", "Benign"], yticklabels=["Malignant", "Benign"])
plt.xlabel("Predicted")
plt.ylabel("Actual")
```

```
plt.title("Confusion Matrix - CatBoost Classifier")
plt.show()
```



In [20]: *#Train an AdaBoost Classifier with different numbers of estimators and compare accuracy*

```
data = load_breast_cancer()
X, y = data.data, data.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

n_estimators_list = [10, 50, 100, 200, 500]

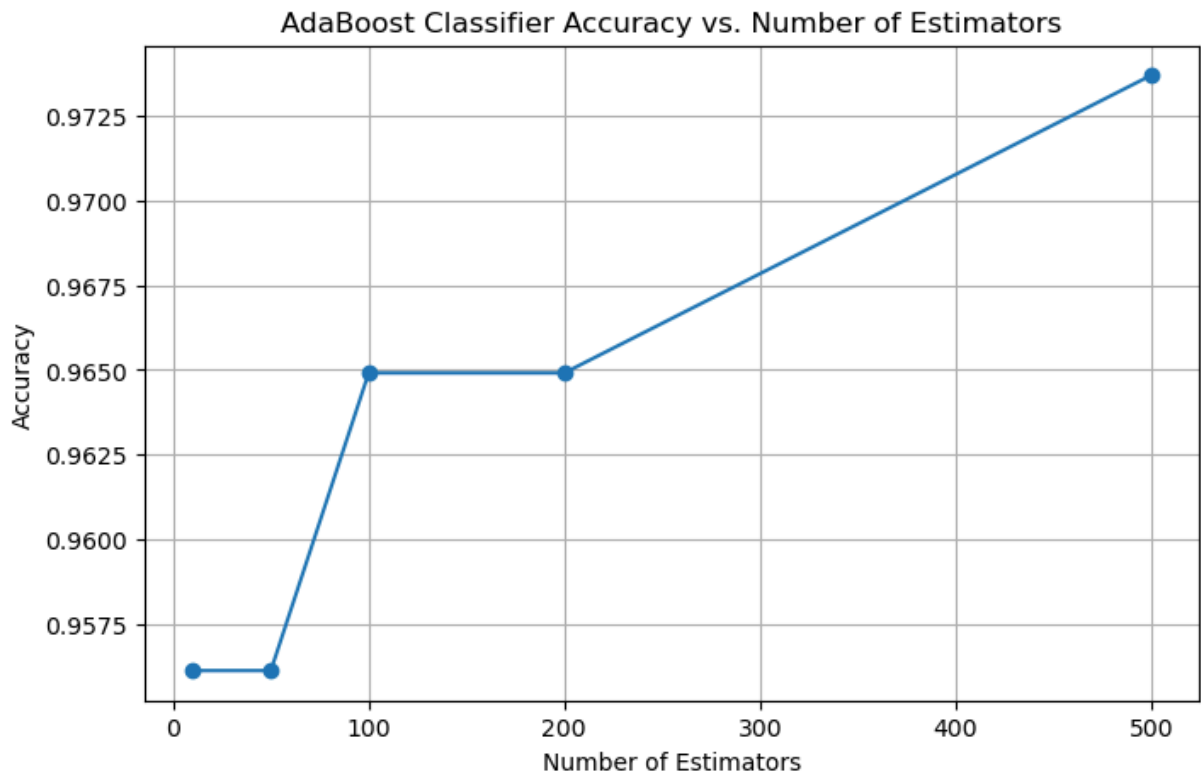
accuracy_scores = []

for n in n_estimators_list:
    ada_clf = AdaBoostClassifier(n_estimators=n, learning_rate=0.1, random_state=42)
    ada_clf.fit(X_train, y_train)
    y_pred = ada_clf.predict(X_test)
    acc = accuracy_score(y_test, y_pred)
    accuracy_scores.append(acc)
    print(f"n_estimators={n}, Accuracy: {acc:.4f}")

plt.figure(figsize=(8, 5))
```

```
plt.plot(n_estimators_list, accuracy_scores, marker='o', linestyle='--')
plt.xlabel("Number of Estimators")
plt.ylabel("Accuracy")
plt.title("AdaBoost Classifier Accuracy vs. Number of Estimators")
plt.grid(True)
plt.show()
```

```
n_estimators=10, Accuracy: 0.9561
n_estimators=50, Accuracy: 0.9561
n_estimators=100, Accuracy: 0.9649
n_estimators=200, Accuracy: 0.9649
n_estimators=500, Accuracy: 0.9737
```



```
In [21]: # Train a Gradient Boosting Classifier and visualize the ROC curve
from sklearn.metrics import roc_curve, auc

data = load_breast_cancer()
X, y = data.data, data.target

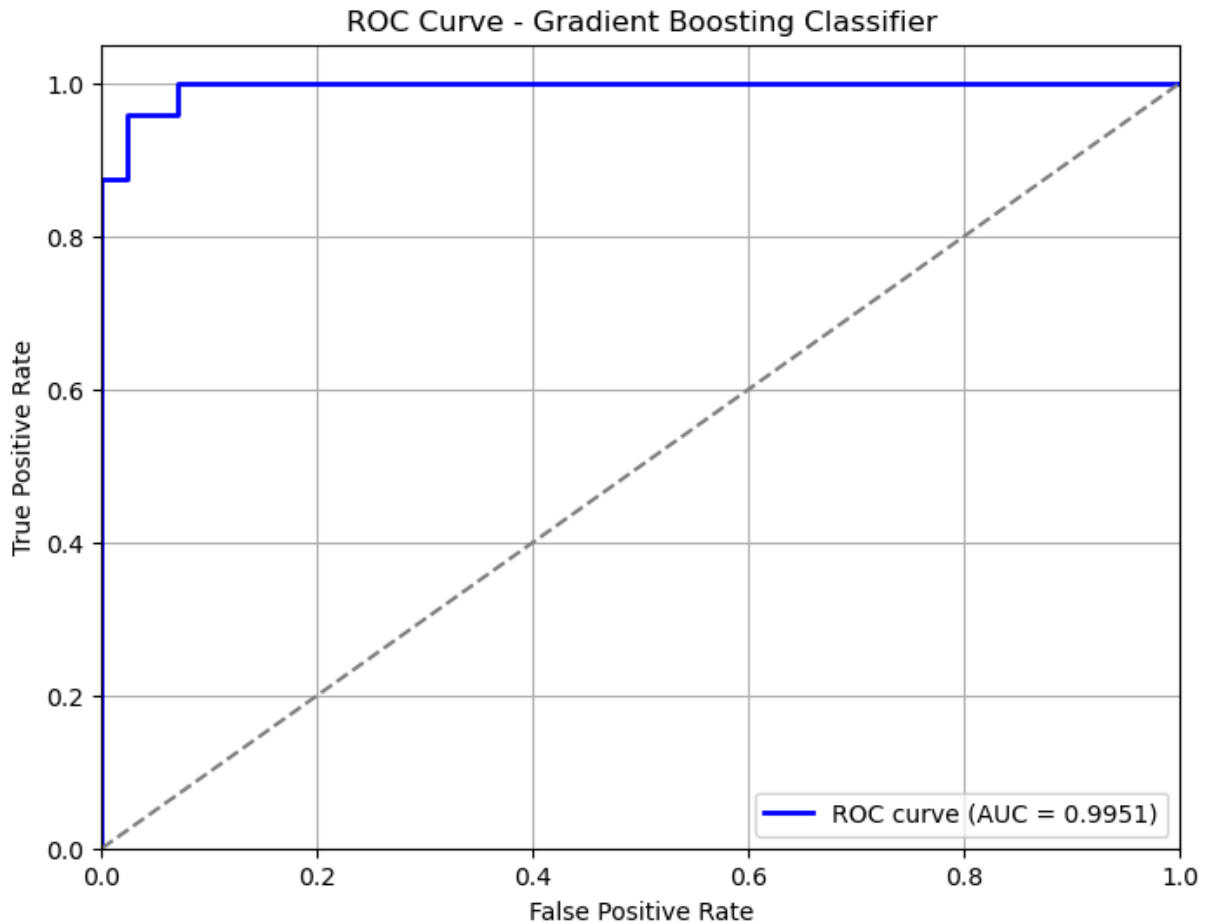
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

gb_clf = GradientBoostingClassifier(n_estimators=100, learning_rate=0.1, random_state=42)
gb_clf.fit(X_train, y_train)

y_scores = gb_clf.predict_proba(X_test)[:, 1]

fpr, tpr, _ = roc_curve(y_test, y_scores)
roc_auc = auc(fpr, tpr)
```

```
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label=f"ROC curve (AUC = {roc_auc:.4f})")
plt.plot([0, 1], [0, 1], color='gray', linestyle='--') # Diagonal line
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve - Gradient Boosting Classifier")
plt.legend(loc="lower right")
plt.grid()
plt.show()
```



```
In [22]: #Train an XGBoost Regressor and tune the Learning rate using GridSearchCV
from xgboost import XGBRegressor
from sklearn.datasets import make_regression
from sklearn.model_selection import GridSearchCV

X, y = make_regression(n_samples=1000, n_features=20, noise=0.1, random_state=42)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```

xgb_reg = XGBRegressor(n_estimators=100, random_state=42)

param_grid = {
    'learning_rate': [0.01, 0.05, 0.1, 0.2, 0.3]
}

grid_search = GridSearchCV(xgb_reg, param_grid, cv=5, scoring='neg_mean_squared_err
grid_search.fit(X_train, y_train)

best_learning_rate = grid_search.best_params_['learning_rate']
print(f"Best Learning Rate: {best_learning_rate}")

best_xgb_reg = XGBRegressor(n_estimators=100, learning_rate=best_learning_rate, ran
best_xgb_reg.fit(X_train, y_train)

y_pred = best_xgb_reg.predict(X_test)

mse = mean_squared_error(y_test, y_pred)
print(f"Mean Squared Error: {mse:.4f}")

```

Best Learning Rate: 0.1

Mean Squared Error: 4845.6934

```

In [24]: #Train a CatBoost Classifier on an imbalanced dataset and compare performance with
from sklearn.metrics import classification_report

X, y = make_classification(n_samples=5000, n_features=20, weights=[0.9, 0.1], random

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_sta

cat_clf_no_weight = CatBoostClassifier(iterations=100, learning_rate=0.1, verbose=0
cat_clf_no_weight.fit(X_train, y_train)
y_pred_no_weight = cat_clf_no_weight.predict(X_test)

class_weights = {0: 1, 1: 10}
cat_clf_weighted = CatBoostClassifier(iterations=100, learning_rate=0.1, class_weig
cat_clf_weighted.fit(X_train, y_train)
y_pred_weighted = cat_clf_weighted.predict(X_test)

print("Without Class Weighting:\n", classification_report(y_test, y_pred_no_weight)
print("With Class Weighting:\n", classification_report(y_test, y_pred_weighted))

fig, axes = plt.subplots(1, 2, figsize=(12, 5))

cm_no_weight = confusion_matrix(y_test, y_pred_no_weight)

```

```

cm_weighted = confusion_matrix(y_test, y_pred_weighted)

sns.heatmap(cm_no_weight, annot=True, fmt="d", cmap="Blues", ax=axes[0])
axes[0].set_title("Confusion Matrix - No Class Weighting")
axes[0].set_xlabel("Predicted")
axes[0].set_ylabel("Actual")

sns.heatmap(cm_weighted, annot=True, fmt="d", cmap="Blues", ax=axes[1])
axes[1].set_title("Confusion Matrix - With Class Weighting")
axes[1].set_xlabel("Predicted")
axes[1].set_ylabel("Actual")

plt.show()

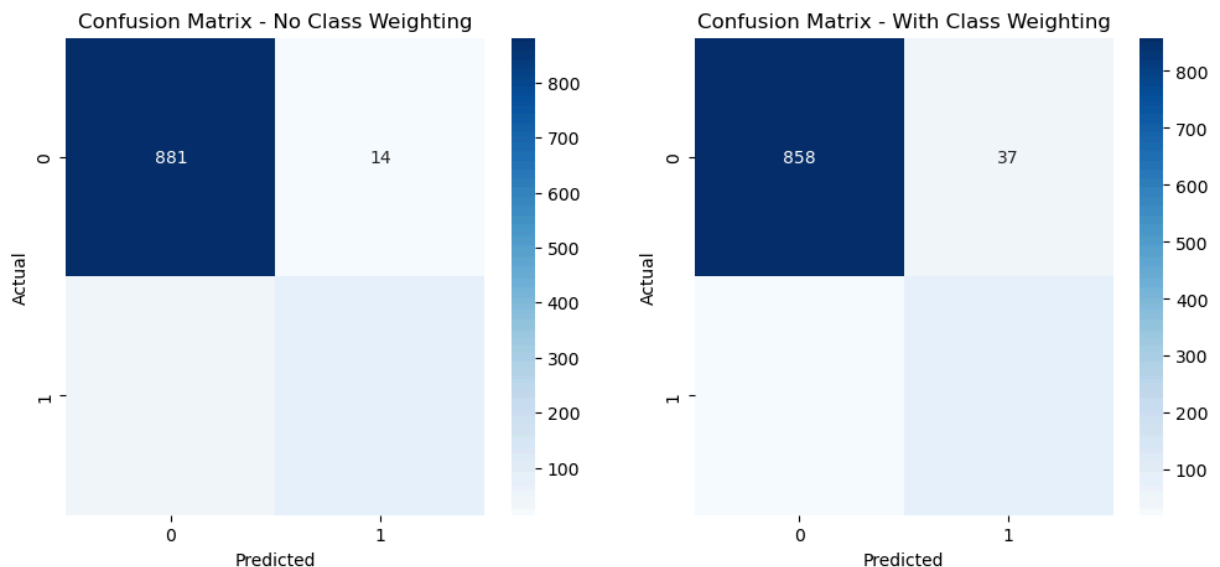
```

Without Class Weighting:

	precision	recall	f1-score	support
0	0.97	0.98	0.98	895
1	0.85	0.73	0.79	105
accuracy			0.96	1000
macro avg	0.91	0.86	0.88	1000
weighted avg	0.96	0.96	0.96	1000

With Class Weighting:

	precision	recall	f1-score	support
0	0.98	0.96	0.97	895
1	0.69	0.80	0.74	105
accuracy			0.94	1000
macro avg	0.84	0.88	0.86	1000
weighted avg	0.95	0.94	0.94	1000



In [25]: *#Train an AdaBoost Classifier and analyze the effect of different Learning rates*

```

data = load_breast_cancer()
X, y = data.data, data.target

```



```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

learning_rates = [0.001, 0.01, 0.1, 0.5, 1.0]

accuracy_scores = []

for lr in learning_rates:
    ada_clf = AdaBoostClassifier(n_estimators=100, learning_rate=lr, random_state=42)
    ada_clf.fit(X_train, y_train)
    y_pred = ada_clf.predict(X_test)
    acc = accuracy_score(y_test, y_pred)
    accuracy_scores.append(acc)
    print(f"Learning Rate={lr}, Accuracy: {acc:.4f}")

plt.figure(figsize=(8, 5))
plt.plot(learning_rates, accuracy_scores, marker='o', linestyle='-', color='b')
plt.xlabel("Learning Rate")
plt.ylabel("Accuracy")
plt.title("Effect of Learning Rate on AdaBoost Classifier Performance")
plt.xscale("log") # Log scale for better visualization
plt.grid(True)
plt.show()

```


Learning Rate=0.001, Accuracy: 0.9386

Learning Rate=0.01, Accuracy: 0.9561

Learning Rate=0.1, Accuracy: 0.9649

Learning Rate=0.5, Accuracy: 0.9737

Learning Rate=1.0, Accuracy: 0.9737



```
In [26]: # Train an XGBoost Classifier for multi-class classification and evaluate using Log
import xgboost as xgb
from sklearn.datasets import load_digits
from sklearn.metrics import log_loss

data = load_digits()
X, y = data.data, data.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

xgb_clf = xgb.XGBClassifier(objective="multi:softprob", num_class=10, n_estimators=100)
xgb_clf.fit(X_train, y_train)

y_pred_proba = xgb_clf.predict_proba(X_test)

logloss = log_loss(y_test, y_pred_proba)
print(f"Log-Loss: {logloss:.4f}")
```

Log-Loss: 0.1732

In []: