

Theoretical

```
In [ ]: #What is K-Nearest Neighbors (KNN) and how does it work?
'''
K-Nearest Neighbors (KNN) is a supervised machine learning algorithm used for class
non-parametric, instance-based (lazy) learning algorithm, meaning it doesn't explic
based on stored training data.

How it works
1.Choose the number of neighbors (K)
K is a hyperparameter that defines how many nearest data points will be considered
2.Calculate the distance
For a given test data point, KNN calculates the distance to every point in the trai
Common distance metrics:
Euclidean Distance (most commonly used)
Manhattan Distance
Minkowski Distance
3.Find the K nearest neighbors
It selects the K closest points from the training dataset
4.Make a prediction
For Classification: The majority class among the K neighbors determines the class o
For Regression: The prediction is the average (or weighted average) of the K neares

Advantages of KNN
Simple and easy to implement.
No need for training; just store the dataset.
Works well with small datasets and low-dimensional data.
'''

#What are the disadvantages of KNN?
'''
Disadvantages of KNN
Computationally expensive for large datasets (since it calculates distances for eve
Sensitive to irrelevant or highly correlated features.
Choosing the right K value can be tricky
'''

# What is the difference between KNN Classification and KNN Regression?
'''
KNN Classification
Output: A category/class label (e.g., "Dog" or "Cat").
Prediction Method:
Finds the K nearest neighbors of the test point.
Uses majority voting—the most common class among the K neighbors is assigned to the
Example Use Case:
Email spam detection (Spam/Not Spam).
Handwriting recognition (Digits 0-9).

KNN Regression
Output: A continuous numerical value (e.g., price, temperature).
Prediction Method:
Finds the K nearest neighbors of the test point.
Computes the average (or weighted average) of their values.
Example Use Case:
Predicting house prices based on nearby houses.
Estimating temperature based on past weather data.
```

```
'''
# What is the role of the distance metric in KNN?
'''

Role of Distance Metric in KNN
The distance metric in K-Nearest Neighbors (KNN) plays a crucial role in determining the model's accuracy. Since KNN makes predictions based on the K nearest neighbors, the way we measure distance affects the model's accuracy and effectiveness of the model.

Why is Distance Important?
It helps in finding the most relevant neighbors for classification or regression.
A poor choice of distance metric can lead to wrong neighbor selection, reducing model accuracy.
It affects the model's sensitivity to different feature scales (e.g., height in cm vs. weight in kg).

Choosing the Right Distance Metric
Euclidean Distance: Best for continuous data with well-scaled features.
Manhattan Distance: Works well for high-dimensional or grid-based data.
Hamming Distance: Used for categorical data.
Cosine Similarity: Ideal for text or high-dimensional sparse data
'''

# What is the Curse of Dimensionality in KNN?
'''

Curse of Dimensionality in KNN
The Curse of Dimensionality refers to the problem where K-Nearest Neighbors (KNN) becomes increasingly inefficient and less accurate as the number of features (dimensions) increases.

Why Does It Happen?
In high-dimensional spaces, data points become more sparse (far apart).
The distance between all points becomes almost the same, making it difficult to distinguish between them.
Since KNN relies on distance measurements, so if all points are equally far, classification becomes unreliable.

Effects of High Dimensionality on KNN
1. Loss of Meaningful Distance
As dimensions increase, Euclidean distance becomes less reliable for finding nearest neighbors.
Most points end up being equidistant, making KNN less effective.
2. Increased Computational Cost
More dimensions mean more distance calculations, leading to longer processing times.
Storing high-dimensional data requires more memory.
3. Overfitting
High-dimensional data often contains irrelevant features (noise).
KNN might consider these irrelevant features, leading to poor generalization on new data.
'''

# How can we choose the best value of K in KNN?
'''

Choosing the right value of K in K-Nearest Neighbors (KNN) is crucial for achieving the best performance. It controls the balance between bias and variance in the model.

Steps to Choose the Best K
1. Try Different Values of K (Grid Search)
Start with small values like K = 1, 3, 5, 7, ... and evaluate performance.
Use a validation set or cross-validation to find the best K.

2. Check for Overfitting & Underfitting
Low K (e.g., K=1, 3):
High variance, sensitive to noise, model may overfit.
Decision boundary is too complex.
High K (e.g., K=15, 20):
Low variance, but may underfit by ignoring local patterns.
Decision boundary is too simple.
```

High bias, model is too smooth and may underfit.
Model ignores local patterns.

3. Use the Elbow Method

Plot K vs. Error Rate (or Accuracy).

Look for the "elbow point" where the error starts to flatten out—this is usually the

4. Consider Odd K for Classification

If the number of classes is 2, an odd value of K prevents ties.

For multi-class problems, choose a value that divides evenly among data.

5. Use Cross-Validation (K-Fold)

Use K-Fold Cross-Validation to test different K values on different subsets of data

This prevents biased selection of K based on a single train-test split.

'''

#What are KD Tree and Ball Tree in KNN?When should you use KD Tree vs. Ball Tree?

'''

In K-Nearest Neighbors (KNN), finding the nearest neighbors requires computing distance for large datasets. To speed up this process, KD-Tree and Ball-Tree are used as efficient search.

KD-Tree (K-Dimensional Tree)

What is it?

A binary tree data structure used for fast nearest neighbor search in low to moderate dimensions. It recursively splits the dataset along the dimension with the highest variance at each node.

How it Works?

Choose a splitting dimension (e.g., x-axis, y-axis).

Select a median value in that dimension as the split point.

Recursively divide points into left (less than median) and right (greater than median) halves.

Perform KNN search efficiently by only checking relevant branches of the tree.

When to Use KD-Tree?

Works well for low-dimensional data (<50 dimensions).

Faster than brute-force search in small-to-medium datasets.

Becomes inefficient in high-dimensional spaces (Curse of Dimensionality).

2. Ball-Tree

What is it?

A hierarchical tree structure used for nearest neighbor search, especially in high-dimensional data. Instead of splitting along a single dimension like KD-Tree, it groups points into clusters (balls) based on proximity.

How it Works?

The dataset is divided into clusters (balls) based on proximity.

Each ball contains a center and a radius that encloses a subset of points.

When searching for nearest neighbors, the tree prunes entire balls that are too far away.

When to Use Ball-Tree?

Better for high-dimensional data (>50 dimensions).

More efficient when Euclidean distance is not ideal.

Slower than KD-Tree in low-dimensional data.

KD-Tree vs. Ball-Tree: Which One to Use?

	Feature	KD Tree	Ball Tree
Data Type	Low to medium dimensions (<50)	High	High
Structure	Splits along dimensions	Good	Good

Efficiency	Fast for small datasets	Better
Distance	Metric Works best with Euclidean Distance	Support

When Not to Use KD-Tree or Ball-Tree?

If the dataset is very high-dimensional (>100 dimensions), both KD-Tree and Ball-Tree are inefficient. In such cases, Approximate Nearest Neighbor (ANN) methods like FAISS, LSH, or HNSW are preferred.

How does feature scaling affect KNN?

...

Feature scaling is crucial in K-Nearest Neighbors (KNN) because KNN relies on distance metrics (like Euclidean or Manhattan distance). If features are not scaled properly, the model may become biased towards features with larger scales.

Why is Feature Scaling Important in KNN?

1. Prevents Dominance of Large-Scale Features

Example:

Feature 1: Age (range: 18-70)

Feature 2: Salary (range: 20,000-200,000)

Since salary values are much larger, they will dominate the distance calculations, leading to poor neighbor selection.

2. Improves Distance Calculation Accuracy

KNN relies on Euclidean Distance (or other metrics), which are sensitive to differences in scale. Proper scaling ensures that all features contribute equally to neighbor selection.

3. Leads to Better Performance & Faster Convergence

Without scaling, KNN might select wrong neighbors, leading to lower accuracy.

Feature scaling makes the algorithm more stable and consistent.

...

#What is PCA (Principal Component Analysis)?

...

Principal Component Analysis (PCA) is a dimensionality reduction technique used to transform high-dimensional data into a lower-dimensional form while preserving the most important patterns.

Why Use PCA?

Reduces Dimensionality - Helps in handling datasets with many features, preventing overfitting.

Removes Redundant Features - Eliminates correlated and less important features.

Improves Computational Efficiency - Makes algorithms like KNN, Logistic Regression, etc., faster.

Better Visualization - Helps in plotting high-dimensional data in 2D or 3D.

...

#How does PCA work?

...

How PCA Works (Step-by-Step)

1. Standardize the Data

Convert features to the same scale (Z-score normalization).

PCA is sensitive to large-scale features, so feature scaling is essential.

2. Compute the Covariance Matrix

Measures how features vary with each other.

Helps find relationships between features.

3. Compute Eigenvalues and Eigenvectors

Eigenvectors represent the directions of new axes (Principal Components).

Eigenvalues represent the amount of variance captured by each component.

4. Select the Top K Principal Components

Sort principal components by eigenvalues (higher eigenvalue = more variance).

Choose the top K components that explain most of the variance (e.g., 95%).

5. Transform Data

Project the original data onto the new PCA space with reduced dimensions.

'''

#What is the geometric intuition behind PCA?

'''

Principal Component Analysis (PCA) can be understood geometrically as a process of capture the variance in the data.

Step-by-Step Geometric Intuition

1. Data as Points in High-Dimensional Space

Imagine your data as a cloud of points in an N-dimensional space (where N is the number of features). Each data point is represented as a vector in this space.

2. Finding the Best New Axes (Principal Components)

PCA rotates the coordinate system to align it with the directions of maximum variance. The new axes (Principal Components) are orthogonal (perpendicular to each other). The first Principal Component (PC1) is the direction where the data varies the most. The second Principal Component (PC2) is perpendicular to PC1 and captures the next

3. Projecting Data onto the New Axes

Once we have the new principal components, we project data onto them. This reduces the number of dimensions while keeping as much variance as possible. Think of it as casting shadows of high-dimensional data onto a lower-dimensional plane.

4. Dropping Less Important Dimensions

PCA removes dimensions with low variance, keeping only the most informative ones. This helps reduce noise and redundancy in the data.

'''

What is the difference between Feature Selection and Feature Extraction?

'''

Feature Selection

Definition: It involves selecting a subset of relevant features (predictors) from the original features themselves.

Purpose: To remove irrelevant, redundant, or less significant features, improving model performance.

Methods:

Filter Methods (e.g., correlation, chi-square test, mutual information)

Wrapper Methods (e.g., recursive feature elimination, forward/backward selection)

Embedded Methods (e.g., LASSO, decision trees)

Example: In a dataset with 100 features, feature selection might keep only 20 of the most relevant ones.

Feature Extraction

Definition: It involves transforming the original features into a new set of features. Transformed features usually have lower dimensionality while preserving essential information.

Purpose: To create a more compact representation of the data while reducing redundancy.

Methods:

Principal Component Analysis (PCA)

Linear Discriminant Analysis (LDA)

Autoencoders (Deep Learning)

t-SNE (for visualization)

Example: Instead of selecting individual features from an image dataset, feature extraction might use PCA to reduce the number of variables that capture the most important variations in the data.

'''

#What are Eigenvalues and Eigenvectors in PCA?

'''

Principal Component Analysis (PCA) is a dimensionality reduction technique that transforms high-dimensional data into a lower-dimensional space while preserving as much variance as possible. The concepts of eigenvalues and eigenvectors are central to PCA.

What are Eigenvalues and Eigenvectors?

Eigenvectors are directions in the data space along which the variance is maximized (principal components) in the transformed space.

Eigenvalues represent the amount of variance carried by each eigenvector. A larger eigenvalue indicates that the corresponding eigenvector captures more information (variance) from the dataset.

How are they used in PCA?

Compute the Covariance Matrix

The covariance matrix captures the relationships (correlations) between features in the dataset.

Find Eigenvalues and Eigenvectors

Solve the characteristic equation

$$(A - \lambda I)v = 0$$

where:

A is the covariance matrix

λ are the eigenvalues

v are the eigenvectors

I is the identity matrix

Sort Eigenvalues and Eigenvectors

The eigenvectors are sorted in descending order based on their eigenvalues. The top k eigenvectors are selected as the principal components.

Select Principal Components

Choose the top

k eigenvectors that explain the most variance (based on eigenvalues) and project the data onto this subspace.

How do you decide the number of components to keep in PCA?

...

Explained Variance (Eigenvalues) - Scree Plot

Concept: Each principal component has an associated eigenvalue, which indicates how much variance it explains. A scree plot shows the eigenvalues for each component, sorted in descending order.

How to use:

Plot a Scree Plot (graph of explained variance vs. number of components).

Identify the "elbow point" where adding more components does not significantly increase the explained variance.

Choose the number of components before the elbow.

Example: If the elbow occurs at 3 components, selecting those 3 is ideal.

2. Retaining a Certain Percentage of Variance

Concept: Retain enough components to capture a high percentage (e.g., 95%) of total variance.

How to use:

Compute cumulative explained variance.

Choose the smallest number of components that retain at least 90-95% of the total variance.

Example: If 5 components explain 95% of the variance, keep those 5.

3. Kaiser Criterion (Eigenvalues > 1)

Concept: Keep only components with eigenvalues greater than 1, as they contribute more than one unit of variance to the data.

How to use:

Compute eigenvalues of principal components.

Keep only those with eigenvalues > 1.

Example: If components 1, 2, and 3 have eigenvalues greater than 1, select those.

4. Cross-Validation with Machine Learning Models

Concept: Use cross-validation to check how PCA components impact model performance.

How to use:

Train a machine learning model (e.g., Logistic Regression, SVM) using different num

Choose the number that gives the best performance (highest accuracy, lowest error).

Example: If 10 components give the best classification accuracy, select those.

5. Parallel Analysis / Randomized Methods

Concept: Compare PCA eigenvalues with those from a randomized dataset of the same s

How to use:

Perform PCA on both real and randomized data.

Keep components where real eigenvalues are higher than random ones.

Example: If the first 7 components have significantly higher eigenvalues than random

'''

Can PCA be used for classification?

'''

No, PCA itself is not a classification algorithm, but it can be used as a preprocessing performance.

How PCA Helps in Classification?

PCA is a dimensionality reduction technique, meaning it transforms the dataset into important patterns. This can be beneficial for classification tasks in the following

1. Removes Redundant or Correlated Features

PCA eliminates correlated features, making it easier for classification models to l

2. Reduces Overfitting

High-dimensional data can lead to overfitting in models like Decision Trees or Neur

PCA reduces dimensionality and prevents models from memorizing noise.

3. Speeds Up Training Time

With fewer features, algorithms like SVM, Logistic Regression, and KNN train faster

4. Handles Multicollinearity

In datasets where independent variables are highly correlated, PCA creates uncorrel

classification models more stable.

'''

#What are the Limitations of PCA?

'''

1. Loss of Interpretability

PCA transforms original features into new principal components, which are linear co

These transformed components are often hard to interpret, making it difficult to ex

applications.

Example:

If a dataset has features like "Age", "Salary", and "Experience", after PCA, the ne

like "Salary Impact."

2. Assumes Linearity

PCA works best when the dataset has a linear structure.

If the data has non-linear relationships, PCA might not capture important patterns.

Alternative: Use kernel PCA (KPCA) or t-SNE for non-linear relationships.

Example:

PCA may fail in recognizing non-linear relationships in image or text classificatio

3. Sensitive to Scaling

PCA is affected by the scale of features since it relies on variance.

Features with large values (e.g., salary in dollars vs. age in years) will dominate

Solution: Always standardize the data (e.g., using `StandardScaler` in `sklearn`).

4. Sensitive to Outliers

Since PCA is based on variance, the presence of outliers can distort the principal

Solution: Detect and remove outliers before applying PCA.

5. May Discard Important Information

PCA prioritizes variance, but low-variance features may still be important for clas

If a classification task depends on a feature with low variance, PCA might remove i

Example:

In fraud detection, rare transactions (low variance) are important, but PCA may dis

6. Computational Cost in High Dimensions

PCA requires calculating eigenvalues and eigenvectors, which can be computationally

Solution: Use Incremental PCA or Randomized PCA for large datasets.

7. Assumes Gaussian Distribution

PCA works best when features follow a Gaussian (normal) distribution.

If data is skewed or multimodal, PCA may not be the best choice.

...

#How do KNN and PCA complement each other?

...

K-Nearest Neighbors (KNN) is a simple, non-parametric classification algorithm that
Euclidean distance). Principal Component Analysis (PCA) is a dimensionality reducti
lower-dimensional space while preserving variance.

When used together, PCA can enhance the performance of KNN in various ways:

1. Reduces Dimensionality to Avoid the Curse of Dimensionality

KNN works poorly in high-dimensional data due to the curse of dimensionality (where
dimensions increase).

PCA reduces the number of features while preserving important information, making K

Example: If a dataset has 500 features, PCA can reduce it to 50 principal componen

2. Speeds Up KNN Computation

KNN is a lazy learner, meaning it stores all training data and computes distances a

With high-dimensional data, distance calculations become expensive.

PCA reduces dimensions, making distance calculations faster.

Example: In image recognition, where each image has thousands of pixels, PCA can re
classification.

3. Reduces Noise and Correlation Among Features

Real-world datasets often have correlated or redundant features that can mislead KN

PCA removes multicollinearity by transforming features into uncorrelated principal

This helps KNN make better distance-based decisions.

Example: In medical diagnosis, multiple lab test results may be correlated. PCA rem
classification accuracy.

4. Prevents Overfitting in High-Dimensional Data

KNN can overfit when there are too many features relative to the number of samples.

PCA reduces dimensions, preventing overfitting while keeping essential variance.

Example: In text classification, converting words into high-dimensional vectors can be used by the model.

5. Handles Collinearity Better

KNN struggles when features are highly correlated.

PCA converts correlated features into independent principal components, making KNN

Example: In stock market prediction, many financial indicators are correlated. PCA handles this better.

#5 How does KNN handle missing values in a dataset?

1. Removing Rows with Missing Values (Simple but Risky)

If the dataset has few missing values, one quick approach is to remove rows with missing values. However, this reduces data size and may lead to bias if missing values are not random.

Use When: The dataset is large, and missing values are rare (<5%).

Avoid If: Missing values are widespread.

2. Imputing Missing Values Using Mean/Median/Mode (Basic Imputation)

Replace missing values with:

Mean (for numerical data)

Median (if data has outliers)

Mode (for categorical data)

Example: If Age = [25, 30, NaN, 35, 40], replace NaN with Mean(Age).

Pros: Easy to implement.

Cons: Ignores relationships between features.

3. KNN-Based Imputation (Best for Preserving Patterns)

KNN can predict missing values by finding the nearest neighbors and averaging their values. This method considers relationships between features.

How It Works:

Identify the k nearest neighbors (based on available features).

Use their values to estimate the missing value.

For numerical data → take the mean of nearest neighbors.

For categorical data → take the mode of nearest neighbors.

Pros: More accurate than simple imputation.

Cons: Computationally expensive for large datasets.

4. Using Machine Learning Models to Predict Missing Values

Train a regression model (for numerical data) or a classification model (for categorical data).

Example: If Age is missing, train a regression model using other features like Income.

Pros: More sophisticated than mean/mode imputation.

Cons: Requires extra computation and model training.

5. Using KNN with Distance-Based Missing Value Handling

Some implementations of KNN can ignore missing values when computing distances.

Instead of discarding a row, KNN calculates distance using only available features.

How It Works:

If feature X is missing for a data point, KNN will calculate distances only based on other features. This prevents data loss but might reduce accuracy.

Pros: Avoids data loss while still using KNN.

Cons: Can lead to biased results if too many features are missing.

#What are the key differences between PCA and Linear Discriminant Analysis (LDA)?

Feature	PCA (Principal Component Analysis)	LDA (Linear Discriminant Analysis)
---------	------------------------------------	------------------------------------

Purpose	Reduces dimensionality by preserving variance	Reduces dim
Type	Unsupervised	Supervise
Works On	Features (independent variables)	Feat
Optimization Criteria	Maximizes variance of data	Maximizes s
Interpretability	Principal components are abstract and may not have clear meaning	Discriminant making t
Use Case	Best for reducing dimensions in datasets without class labels	Best for cl
Computational Cost	Moderate (Eigenvalue decomposition)	Highe
Handles Multicollinearity?	Yes (removes correlation between features)	
Output Components	At most $\min(n_{\text{samples}}, n_{\text{features}})$	
Effect on Data Distribution	Retains total variance but does not guarantee separation between classes	

Practical

```
In [1]: #Train a KNN Classifier on the Iris dataset and print model accuracy
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score
from sklearn.datasets import load_iris

iris = load_iris()
X = iris.data
y = iris.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_sta

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train_scaled, y_train)

y_pred = knn.predict(X_test_scaled)

accuracy = accuracy_score(y_test, y_pred)
print(f"KNN Model Accuracy on Iris Dataset: {accuracy:.4f}")
```

KNN Model Accuracy on Iris Dataset: 1.0000

```

In [3]: #Train a KNN Regressor on a synthetic dataset and evaluate using Mean Squared Error
import matplotlib.pyplot as plt
from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import mean_squared_error

# Generate synthetic dataset
np.random.seed(42)
X = np.random.rand(100, 1) * 10 # Features (100 samples, 1 feature)
y = np.sin(X).ravel() + np.random.normal(0, 0.3, X.shape[0]) # Target with noise

# Split dataset
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardize features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Train KNN Regressor
knn_regressor = KNeighborsRegressor(n_neighbors=5)
knn_regressor.fit(X_train_scaled, y_train)

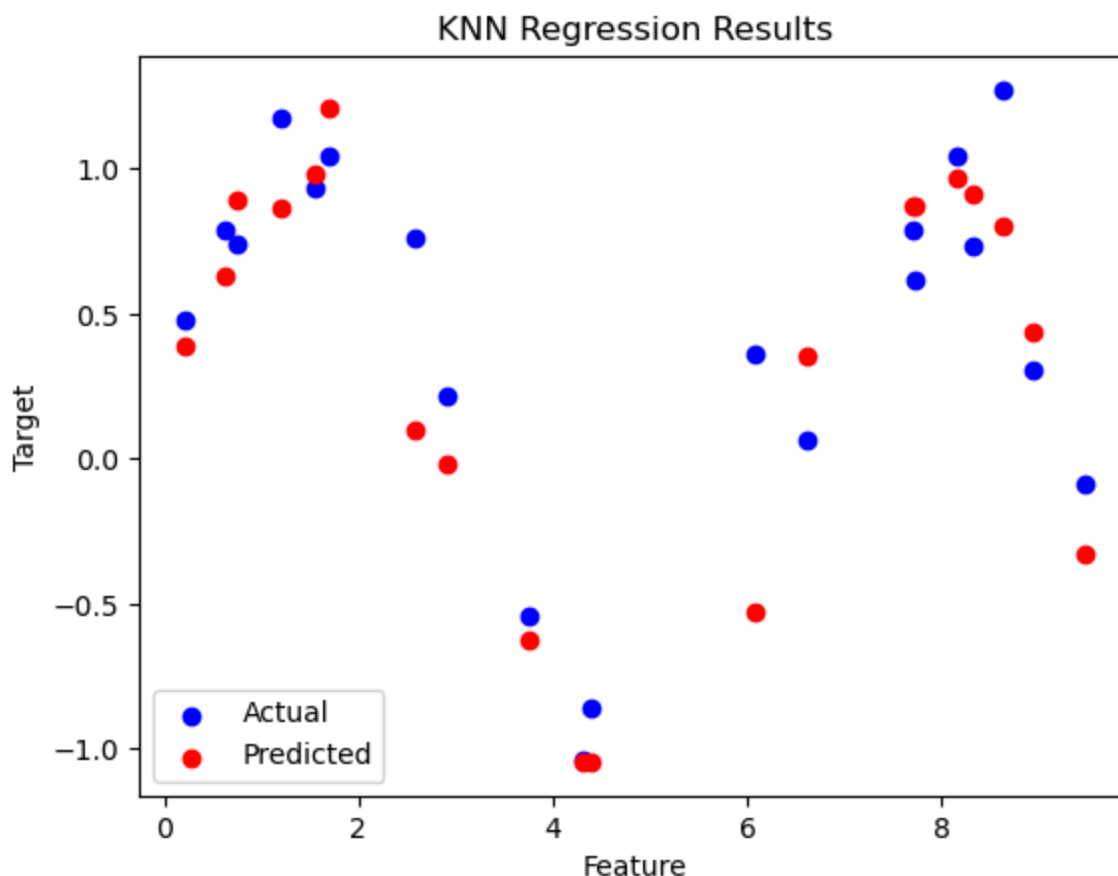
# Predict
y_pred = knn_regressor.predict(X_test_scaled)

# Evaluate using MSE
mse = mean_squared_error(y_test, y_pred)
print(f"KNN Regressor MSE: {mse:.4f}")

# Plot results
plt.scatter(X_test, y_test, color="blue", label="Actual")
plt.scatter(X_test, y_pred, color="red", label="Predicted")
plt.xlabel("Feature")
plt.ylabel("Target")
plt.legend()
plt.title("KNN Regression Results")
plt.show()

```

KNN Regressor MSE: 0.0997



In [4]: *#Train a KNN Classifier using different distance metrics (Euclidean and Manhattan)*

```
iris = load_iris()
X = iris.data
y = iris.target

# Split dataset
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardize features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Train KNN with Euclidean distance
knn_euclidean = KNeighborsClassifier(n_neighbors=5, metric='euclidean')
knn_euclidean.fit(X_train_scaled, y_train)
y_pred_euclidean = knn_euclidean.predict(X_test_scaled)
accuracy_euclidean = accuracy_score(y_test, y_pred_euclidean)

# Train KNN with Manhattan distance
knn_manhattan = KNeighborsClassifier(n_neighbors=5, metric='manhattan')
knn_manhattan.fit(X_train_scaled, y_train)
y_pred_manhattan = knn_manhattan.predict(X_test_scaled)
accuracy_manhattan = accuracy_score(y_test, y_pred_manhattan)

# Print results
```

```
print(f"KNN Accuracy (Euclidean Distance): {accuracy_euclidean:.4f}")
print(f"KNN Accuracy (Manhattan Distance): {accuracy_manhattan:.4f}")
```

KNN Accuracy (Euclidean Distance): 1.0000

KNN Accuracy (Manhattan Distance): 1.0000

In [8]: *#Train a KNN Classifier with different values of K and visualize decision boundaries*

```
from sklearn.datasets import make_classification

from matplotlib.colors import ListedColormap



# Generate synthetic dataset (Fix: Set n_redundant=0)
X, y = make_classification(n_samples=200, n_features=2, n_informative=2, n_redundant=0,
                           n_classes=2, n_clusters_per_class=1, random_state=42)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


# Function to plot decision boundaries
def plot_decision_boundary(X, y, model, title):
    h = 0.02
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))

    Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)

    plt.contourf(xx, yy, Z, alpha=0.3, cmap=ListedColormap(['#FFAAAA', '#AAAAFF']))
    plt.scatter(X[:, 0], X[:, 1], c=y, edgecolors='k', cmap=ListedColormap(['#FF0000', '#0000FF']))
    plt.title(title)
    plt.show()

# Train and visualize KNN for different values of K
k_values = [1, 5, 15]
for k in k_values:
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train, y_train)
    plot_decision_boundary(X, y, knn, f"KNN Decision Boundary (k={k})")
```





In []: *#Apply Feature Scaling before training a KNN model and compare results with unscale*

```
# Load dataset
iris = load_iris()
X = iris.data
y = iris.target

# Split dataset
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_sta

# Train KNN without scaling
knn_unscaled = KNeighborsClassifier(n_neighbors=5)
knn_unscaled.fit(X_train, y_train)
y_pred_unscaled = knn_unscaled.predict(X_test)
accuracy_unscaled = accuracy_score(y_test, y_pred_unscaled)

# Apply feature scaling
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Train KNN with scaling
knn_scaled = KNeighborsClassifier(n_neighbors=5)
knn_scaled.fit(X_train_scaled, y_train)
y_pred_scaled = knn_scaled.predict(X_test_scaled)
accuracy_scaled = accuracy_score(y_test, y_pred_scaled)

# Print comparison results
```

```
print(f"KNN Accuracy without Scaling: {accuracy_unscaled:.4f}")
print(f"KNN Accuracy with Scaling: {accuracy_scaled:.4f}")
```

```
In [10]: #Train a PCA model on synthetic data and print the explained variance ratio for each component
import numpy as np
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA

# Generate synthetic data (100 samples, 5 features)
np.random.seed(42)
X = np.random.rand(100, 5) * 10 # Random values between 0 and 10

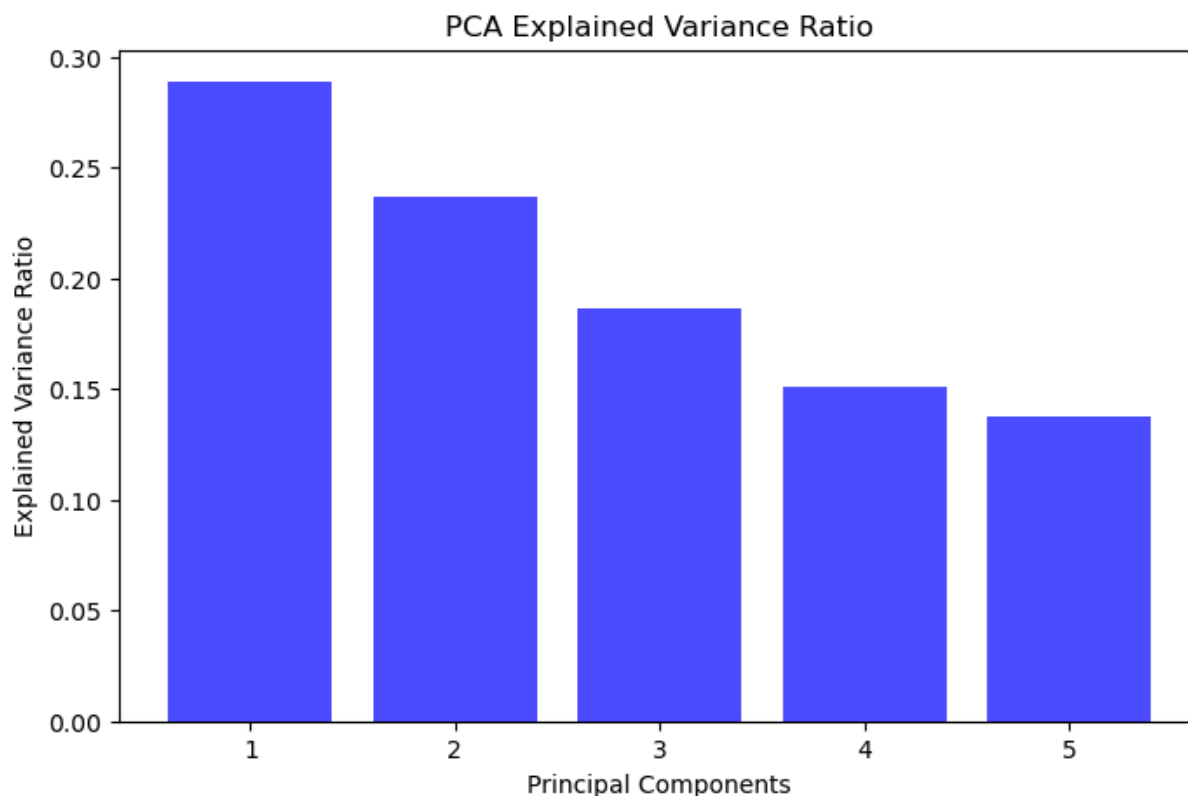
# Apply PCA
pca = PCA(n_components=5) # Keep all components
pca.fit(X)

# Print explained variance ratio
print("Explained Variance Ratio for each Principal Component:")
for i, ratio in enumerate(pca.explained_variance_ratio_):
    print(f"Component {i+1}: {ratio:.4f}")

# Plot the explained variance ratio
plt.figure(figsize=(8, 5))
plt.bar(range(1, 6), pca.explained_variance_ratio_, alpha=0.7, color='b')
plt.xlabel('Principal Components')
plt.ylabel('Explained Variance Ratio')
plt.title('PCA Explained Variance Ratio')
plt.show()
```

Explained Variance Ratio for each Principal Component:

Component 1: 0.2886
 Component 2: 0.2364
 Component 3: 0.1862
 Component 4: 0.1510
 Component 5: 0.1379



```
In [11]: #Apply PCA before training a KNN Classifier and compare accuracy with and without PCA
import numpy as np
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.datasets import make_classification

# Step 1: Generate synthetic dataset
X, y = make_classification(n_samples=500, n_features=10, n_informative=8, n_classes=2)

# Step 2: Split into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Step 3: Train KNN classifier without PCA
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)
accuracy_without_pca = accuracy_score(y_test, y_pred)

# Step 4: Apply PCA (reduce to 5 components)
pca = PCA(n_components=5)
X_train_pca = pca.fit_transform(X_train)
X_test_pca = pca.transform(X_test)


# Train KNN classifier on PCA-transformed data
knn_pca = KNeighborsClassifier(n_neighbors=5)
knn_pca.fit(X_train_pca, y_train)
y_pred_pca = knn_pca.predict(X_test_pca)
accuracy_with_pca = accuracy_score(y_test, y_pred_pca)
```

```
# Step 5: Compare accuracy
print(f"Accuracy without PCA: {accuracy_without_pca:.4f}")
print(f"Accuracy with PCA: {accuracy_with_pca:.4f}")

# Plot explained variance ratio
plt.figure(figsize=(8, 5))
plt.bar(range(1, 6), pca.explained_variance_ratio_, alpha=0.7, color='b')
plt.xlabel('Principal Components')
plt.ylabel('Explained Variance Ratio')
plt.title('PCA Explained Variance Ratio')
plt.show()
```

Accuracy without PCA: 0.7500

Accuracy with PCA: 0.6600



```
In [12]: # Perform Hyperparameter Tuning on a KNN Classifier using GridSearchCV
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import GridSearchCV

# Step 1: Generate a synthetic dataset
X, y = make_classification(n_samples=500, n_features=10, n_informative=8, n_classes=2)

# Step 2: Split into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Step 3: Define KNN model
knn = KNeighborsClassifier()

# Step 4: Define the hyperparameter grid
param_grid = {
    'n_neighbors': [3, 5, 7, 9, 11], # Number of neighbors
```

```

    'weights': ['uniform', 'distance'], # Weighting method
    'metric': ['euclidean', 'manhattan', 'minkowski'] # Distance metric
}

# Step 5: Apply GridSearchCV
grid_search = GridSearchCV(knn, param_grid, cv=5, scoring='accuracy', n_jobs=-1)
grid_search.fit(X_train, y_train)

# Step 6: Get best parameters and accuracy
best_params = grid_search.best_params_
best_model = grid_search.best_estimator_

# Step 7: Evaluate on the test set
y_pred = best_model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)

# Step 8: Print results
print(f"Best Hyperparameters: {best_params}")
print(f"Test Set Accuracy with Best Hyperparameters: {accuracy:.4f}")

```

Best Hyperparameters: {'metric': 'euclidean', 'n_neighbors': 5, 'weights': 'distance'}

Test Set Accuracy with Best Hyperparameters: 0.7400

In [13]: #Train a KNN Classifier and check the number of misclassified samples

```

# Step 1: Generate synthetic dataset
X, y = make_classification(n_samples=500, n_features=10, n_informative=8, n_classes=2)

# Step 2: Split into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Step 3: Train KNN Classifier (k=5)
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train, y_train)

# Step 4: Make predictions
y_pred = knn.predict(X_test)

# Step 5: Compute accuracy
accuracy = accuracy_score(y_test, y_pred)

# Step 6: Calculate number of misclassified samples
misclassified_samples = (y_test != y_pred).sum()

# Step 7: Print results
print(f"Test Set Accuracy: {accuracy:.4f}")
print(f"Number of Misclassified Samples: {misclassified_samples} out of {len(y_test)}")

```

Test Set Accuracy: 0.7500

Number of Misclassified Samples: 25 out of 100

In [14]: #Train a PCA model and visualize the cumulative explained variance.

```

from sklearn.decomposition import PCA

```

```


# Step 1: Generate synthetic dataset
X, y = make_classification(n_samples=500, n_features=10, n_informative=8, n_classes

# Step 2: Apply PCA (keep all components)
pca = PCA()
pca.fit(X)

# Step 3: Compute cumulative explained variance
cumulative_variance = np.cumsum(pca.explained_variance_ratio_)

# Step 4: Plot cumulative explained variance
plt.figure(figsize=(8, 5))
plt.plot(range(1, len(cumulative_variance) + 1), cumulative_variance, marker='o', l
plt.xlabel('Number of Principal Components')
plt.ylabel('Cumulative Explained Variance')
plt.title('PCA Cumulative Explained Variance')
plt.grid(True)
plt.axhline(y=0.95, color='r', linestyle='--', label="95% Variance Threshold")
plt.legend()
plt.show()

```



In [15]: #Train a KNN Classifier using different values of the weights parameter (uniform vs

```

# Step 1: Generate a synthetic dataset
X, y = make_classification(n_samples=500, n_features=10, n_informative=8, n_classes

# Step 2: Split into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_sta

```

```

# Step 3: Train KNN classifier with "uniform" weights
knn_uniform = KNeighborsClassifier(n_neighbors=5, weights='uniform')
knn_uniform.fit(X_train, y_train)
y_pred_uniform = knn_uniform.predict(X_test)
accuracy_uniform = accuracy_score(y_test, y_pred_uniform)

# Step 4: Train KNN classifier with "distance" weights
knn_distance = KNeighborsClassifier(n_neighbors=5, weights='distance')
knn_distance.fit(X_train, y_train)
y_pred_distance = knn_distance.predict(X_test)
accuracy_distance = accuracy_score(y_test, y_pred_distance)

# Step 5: Compare accuracy
print(f"Accuracy with Uniform Weights: {accuracy_uniform:.4f}")
print(f"Accuracy with Distance Weights: {accuracy_distance:.4f}")

```

Accuracy with Uniform Weights: 0.7500

Accuracy with Distance Weights: 0.7400

In [16]: # Train a KNN Regressor and analyze the effect of different K values on performance

```

from sklearn.neighbors import KNeighborsRegressor
from sklearn.datasets import make_regression
from sklearn.metrics import mean_squared_error

# Step 1: Generate synthetic regression data
X, y = make_regression(n_samples=500, n_features=1, noise=15, random_state=42)

# Step 2: Split into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

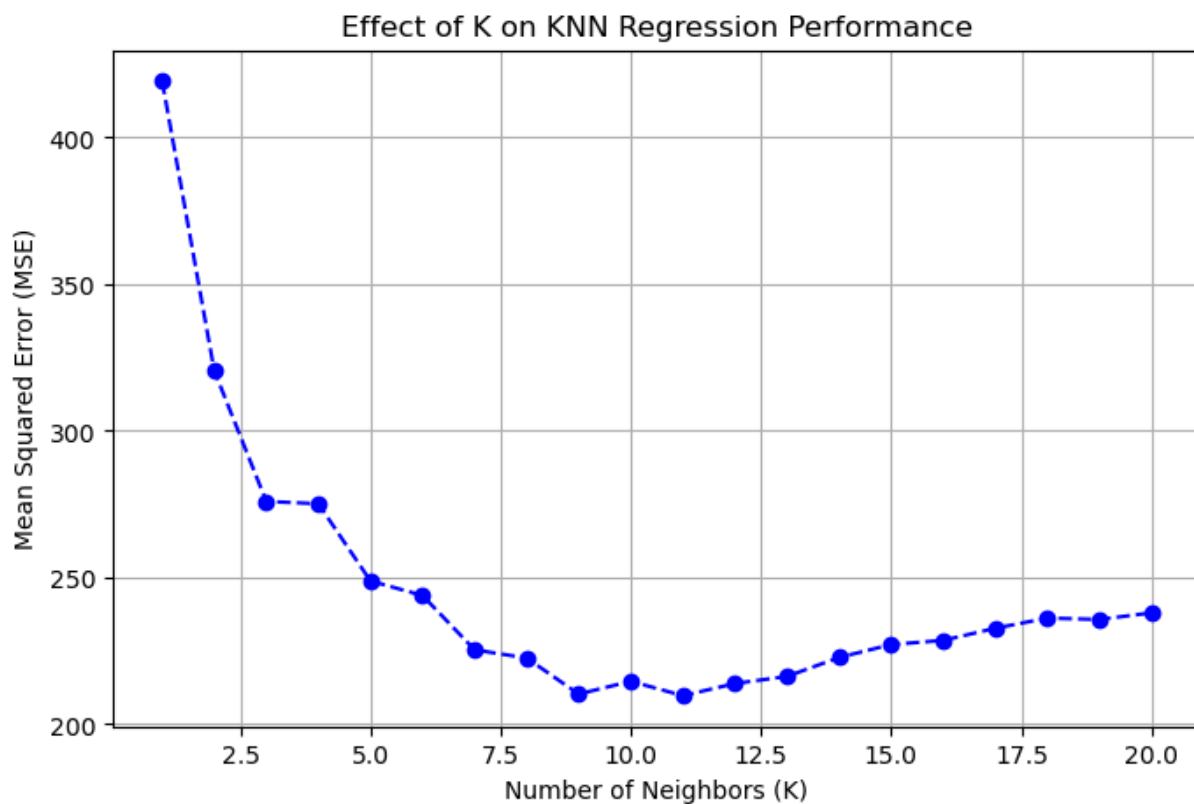
# Step 3: Train KNN Regressors with different values of K
k_values = range(1, 21) # Testing K values from 1 to 20
mse_values = []

for k in k_values:
    knn_reg = KNeighborsRegressor(n_neighbors=k)
    knn_reg.fit(X_train, y_train)
    y_pred = knn_reg.predict(X_test)
    mse = mean_squared_error(y_test, y_pred)
    mse_values.append(mse)

# Step 4: Plot K vs. MSE
plt.figure(figsize=(8, 5))
plt.plot(k_values, mse_values, marker='o', linestyle='--', color='b')
plt.xlabel("Number of Neighbors (K)")
plt.ylabel("Mean Squared Error (MSE)")
plt.title("Effect of K on KNN Regression Performance")
plt.grid(True)
plt.show()

# Step 5: Print best K value
best_k = k_values[np.argmin(mse_values)]
print(f"Best K value: {best_k} (Lowest MSE: {min(mse_values):.4f})")

```



Best K value: 11 (Lowest MSE: 209.7065)

```
In [17]: # Implement KNN Imputation for handling missing values in a dataset

from sklearn.impute import KNNImputer

# Step 1: Create a synthetic dataset with missing values
np.random.seed(42)
data = np.random.rand(10, 5) * 10 # 10 samples, 5 features
data[2, 1] = np.nan # Introduce missing value
data[5, 3] = np.nan
data[7, 4] = np.nan

df = pd.DataFrame(data, columns=[f'Feature_{i+1}' for i in range(5)])
print("Original Data with Missing Values:\n", df)

# Step 2: Apply KNN Imputation
knn_imputer = KNNImputer(n_neighbors=3) # Using K=3 for imputation
imputed_data = knn_imputer.fit_transform(df)

# Step 3: Convert back to DataFrame
df_imputed = pd.DataFrame(imputed_data, columns=df.columns)
print("\nData After KNN Imputation:\n", df_imputed)
```

Original Data with Missing Values:

	Feature_1	Feature_2	Feature_3	Feature_4	Feature_5
0	3.745401	9.507143	7.319939	5.986585	1.560186
1	1.559945	0.580836	8.661761	6.011150	7.080726
2	0.205845	NaN	8.324426	2.123391	1.818250
3	1.834045	3.042422	5.247564	4.319450	2.912291
4	6.118529	1.394939	2.921446	3.663618	4.560700
5	7.851760	1.996738	5.142344	NaN	0.464504
6	6.075449	1.705241	0.650516	9.488855	9.656320
7	8.083973	3.046138	0.976721	6.842330	NaN
8	1.220382	4.951769	0.343885	9.093204	2.587800
9	6.625223	3.117111	5.200680	5.467103	1.848545


Data After KNN Imputation:

	Feature_1	Feature_2	Feature_3	Feature_4	Feature_5
0	3.745401	9.507143	7.319939	5.986585	1.560186
1	1.559945	0.580836	8.661761	6.011150	7.080726
2	0.205845	4.376801	8.324426	2.123391	1.818250
3	1.834045	3.042422	5.247564	4.319450	2.912291
4	6.118529	1.394939	2.921446	3.663618	4.560700
5	7.851760	1.996738	5.142344	5.324350	0.464504
6	6.075449	1.705241	0.650516	9.488855	9.656320
7	8.083973	3.046138	0.976721	6.842330	5.355188
8	1.220382	4.951769	0.343885	9.093204	2.587800
9	6.625223	3.117111	5.200680	5.467103	1.848545

```
In [18]: #Train a PCA model and visualize the data projection onto the first two principal c
X, y = make_classification(n_samples=500, n_features=10, n_informative=8, n_classes

pca = PCA(n_components=2)
X_pca = pca.fit_transform(X)

plt.figure(figsize=(8, 6))
scatter = plt.scatter(X_pca[:, 0], X_pca[:, 1], c=y, cmap='viridis', edgecolors='k'
plt.xlabel("Principal Component 1")
plt.ylabel("Principal Component 2")
plt.title("Data Projection onto First Two Principal Components")
plt.colorbar(label="Class Label")
plt.grid(True)
plt.show()
```



In [19]: *# Train a KNN Classifier using the KD Tree and Ball Tree algorithms and compare per*

```
import time
```

```
# Step 1: Generate synthetic dataset
```

```
X, y = make_classification(n_samples=1000, n_features=10, n_informative=8, n_classes=3)
```

```
# Step 2: Split into training and testing sets (80% train, 20% test)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Step 3: Train KNN Classifier using KD Tree
```

```
start_time = time.time()
```

```
knn_kd = KNeighborsClassifier(n_neighbors=5, algorithm='kd_tree')
```

```
knn_kd.fit(X_train, y_train)
```

```
y_pred_kd = knn_kd.predict(X_test)
```

```
kd_time = time.time() - start_time
```

```
kd_accuracy = accuracy_score(y_test, y_pred_kd)
```

```
# Step 4: Train KNN Classifier using Ball Tree
```

```
start_time = time.time()
```

```
knn_ball = KNeighborsClassifier(n_neighbors=5, algorithm='ball_tree')
```

```
knn_ball.fit(X_train, y_train)
```

```
y_pred_ball = knn_ball.predict(X_test)
```

```
ball_time = time.time() - start_time
```

```
ball_accuracy = accuracy_score(y_test, y_pred_ball)
```



```
# Step 5: Compare Performance
```

```
print(f"KD Tree - Accuracy: {kd_accuracy:.4f}, Training Time: {kd_time:.4f} sec")
print(f"Ball Tree - Accuracy: {ball_accuracy:.4f}, Training Time: {ball_time:.4f} s
```

KD Tree - Accuracy: 0.9000, Training Time: 0.0143 sec

Ball Tree - Accuracy: 0.9000, Training Time: 0.0178 sec

In [20]: # Train a PCA model on a high-dimensional dataset and visualize the Scree plot

```
# Step 1: Generate a high-dimensional dataset (500 samples, 50 features)
```

```
X, y = make_classification(n_samples=500, n_features=50, n_informative=30, random_s
```

```
# Step 2: Apply PCA (keeping all components)
```

```
pca = PCA()
```

```
pca.fit(X)
```

```
# Step 3: Explained variance ratio (Scree plot)
```

```
plt.figure(figsize=(8, 5))
```

```
plt.plot(range(1, len(pca.explained_variance_ratio_) + 1), pca.explained_variance_r
```

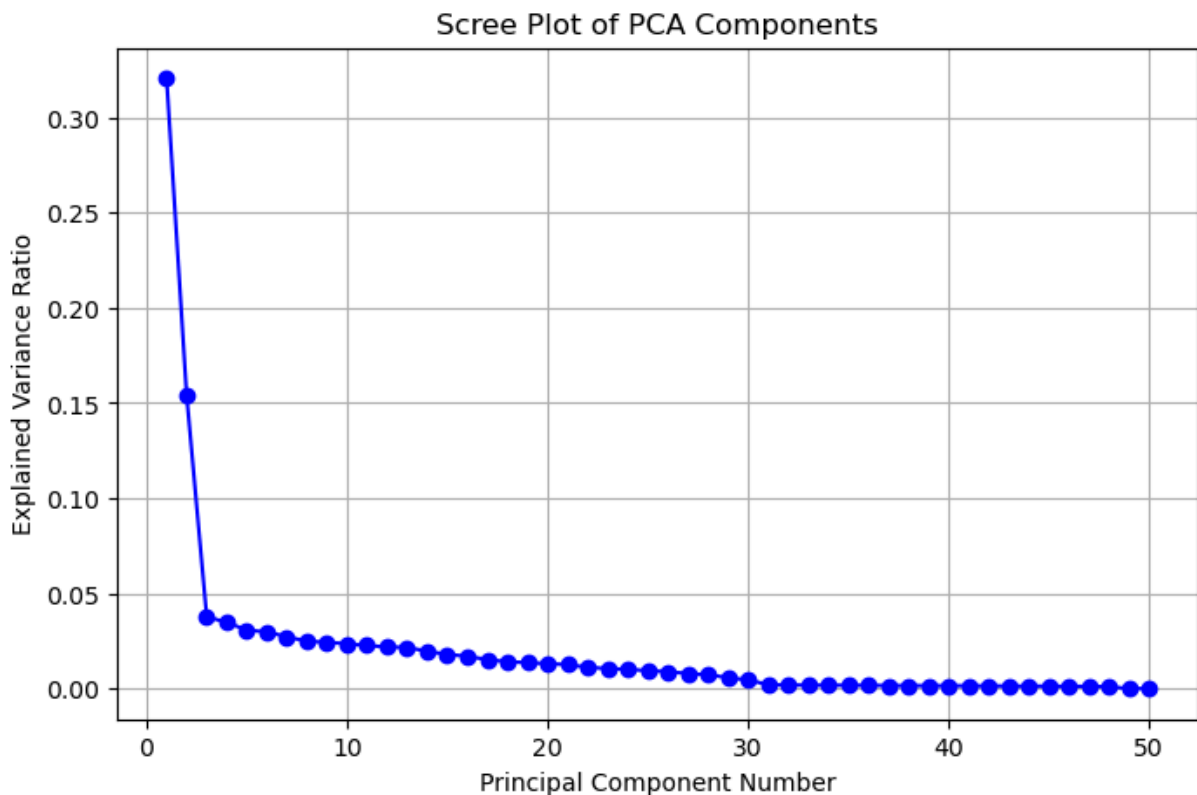
```
plt.xlabel("Principal Component Number")
```

```
plt.ylabel("Explained Variance Ratio")
```

```
plt.title("Scree Plot of PCA Components")
```

```
plt.grid(True)
```

```
plt.show()
```



In [21]: # Train a KNN Classifier and evaluate performance using Precision, Recall, and F1-S

```
from sklearn.metrics import precision_score, recall_score, f1_score, classification
```

```
# Step 1: Generate synthetic dataset (multi-class classification)
```

```
X, y = make_classification(n_samples=1000, n_features=10, n_informative=8, n_classe
```

```

# Step 2: Split into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_sta

# Step 3: Train KNN Classifier
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)

# Step 4: Evaluate Performance
precision = precision_score(y_test, y_pred, average='weighted')
recall = recall_score(y_test, y_pred, average='weighted')
f1 = f1_score(y_test, y_pred, average='weighted')

print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1-Score: {f1:.4f}")

# Print detailed classification report
print("\nClassification Report:\n", classification_report(y_test, y_pred))

```

Precision: 0.9001

Recall: 0.9000

F1-Score: 0.8999

Classification Report:

	precision	recall	f1-score	support
0	0.90	0.93	0.91	67
1	0.91	0.89	0.90	70
2	0.89	0.89	0.89	63
accuracy			0.90	200
macro avg	0.90	0.90	0.90	200
weighted avg	0.90	0.90	0.90	200

In [22]: # Train a PCA model and analyze the effect of different numbers of components on ac

```

# Step 1: Generate synthetic classification data
X, y = make_classification(n_samples=1000, n_features=20, n_informative=15, random_

# Step 2: Split into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_sta

# Step 3: Test different numbers of PCA components
components_list = range(1, 21) # Try PCA with 1 to 20 components
accuracy_scores = []

for n_components in components_list:
    # Apply PCA
    pca = PCA(n_components=n_components)
    X_train_pca = pca.fit_transform(X_train)
    X_test_pca = pca.transform(X_test)

```

```

# Train KNN Classifier
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train_pca, y_train)
y_pred = knn.predict(X_test_pca)

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
accuracy_scores.append(accuracy)

# Step 4: Plot Accuracy vs. Number of Components
plt.figure(figsize=(8, 5))
plt.plot(components_list, accuracy_scores, marker='o', linestyle='--', color='b')
plt.xlabel("Number of Principal Components")
plt.ylabel("Accuracy")
plt.title("Effect of PCA Components on KNN Accuracy")
plt.grid(True)
plt.show()


# Step 5: Print best number of components
best_components = components_list[np.argmax(accuracy_scores)]
print(f"Best number of PCA components: {best_components} (Highest Accuracy: {max(ac

```

```

/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-packages/joblib/external
s/loky/backend/context.py:110: UserWarning: Could not find the number of physical co
res for the following reason:
found 0 physical cores < 1
Returning the number of logical cores instead. You can silence this warning by setti
ng LOKY_MAX_CPU_COUNT to the number of cores you want to use.
  warnings.warn(
    File "/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-packages/joblib/e
xternals/loky/backend/context.py", line 217, in _count_physical_cores
      raise ValueError(

```



Best number of PCA components: 16 (Highest Accuracy: 0.9650)

```
In [23]: #Train a KNN Classifier with different leaf_size values and compare accuracy

# Step 1: Generate synthetic classification dataset
X, y = make_classification(n_samples=1000, n_features=20, n_informative=15, random_

# Step 2: Split into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_sta

# Step 3: Train KNN with different leaf_size values
leaf_sizes = [5, 10, 20, 30, 50, 100] # Different values to test
accuracy_scores = []

for leaf in leaf_sizes:
    knn = KNeighborsClassifier(n_neighbors=5, leaf_size=leaf) # Keeping k=5 consta
    knn.fit(X_train, y_train)
    y_pred = knn.predict(X_test)

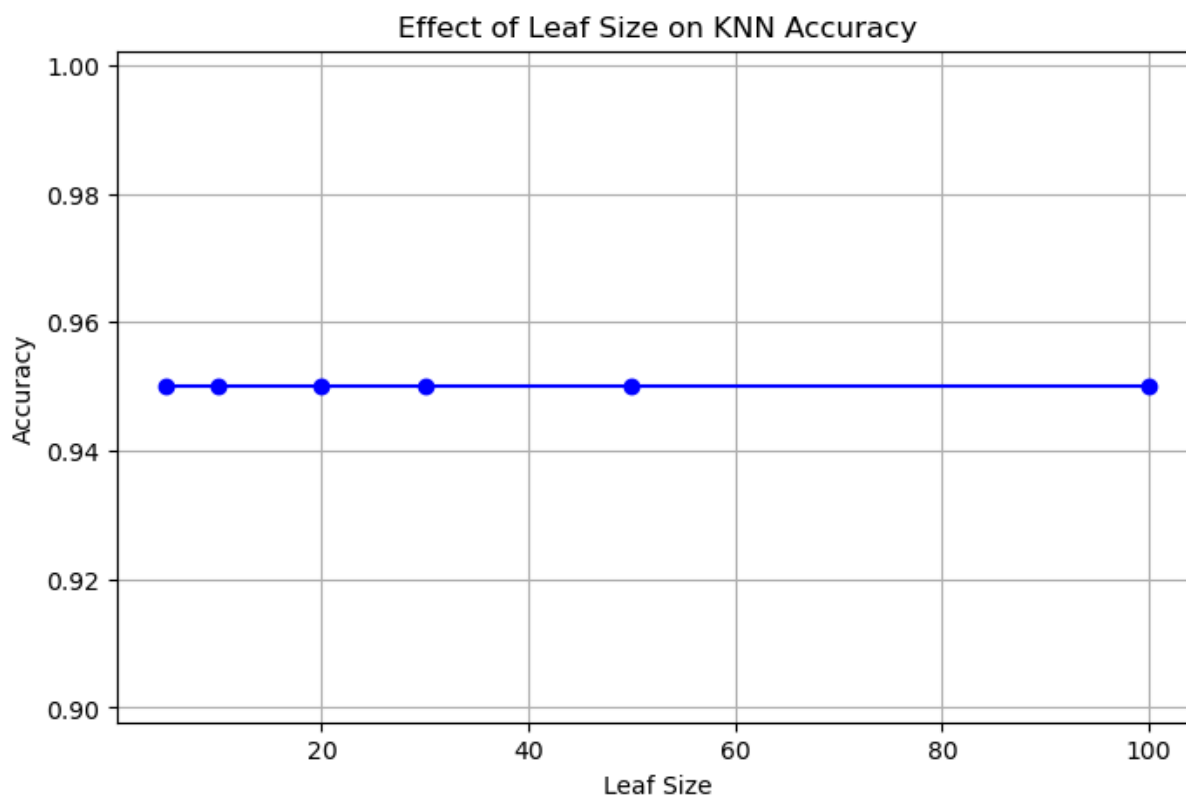
    accuracy = accuracy_score(y_test, y_pred)
    accuracy_scores.append(accuracy)
    print(f"Leaf Size {leaf} -> Accuracy: {accuracy:.4f}")

# Step 4: Plot Accuracy vs. Leaf Size
plt.figure(figsize=(8, 5))
plt.plot(leaf_sizes, accuracy_scores, marker='o', linestyle='--', color='b')
plt.xlabel("Leaf Size")
plt.ylabel("Accuracy")
plt.title("Effect of Leaf Size on KNN Accuracy")
plt.grid(True)
```

```
plt.show()

# Step 5: Print the best leaf size
best_leaf_size = leaf_sizes[np.argmax(accuracy_scores)]
print(f"\nBest leaf_size: {best_leaf_size} (Highest Accuracy: {max(accuracy_scores)})
```

Leaf Size 5 -> Accuracy: 0.9500
 Leaf Size 10 -> Accuracy: 0.9500
 Leaf Size 20 -> Accuracy: 0.9500
 Leaf Size 30 -> Accuracy: 0.9500
 Leaf Size 50 -> Accuracy: 0.9500
 Leaf Size 100 -> Accuracy: 0.9500



Best leaf_size: 5 (Highest Accuracy: 0.9500)

```
In [25]: #Train a PCA model and visualize how data points are transformed before and after P
from sklearn.datasets import make_classification

# Step 1: Generate synthetic dataset (3 features for visualization)
X, y = make_classification(n_samples=500, n_features=10, n_informative=6, n_redunda

# Step 2: Scatter Plot of Original 3D Data
fig = plt.figure(figsize=(12, 5))


ax = fig.add_subplot(121, projection='3d')
scatter = ax.scatter(X[:, 0], X[:, 1], X[:, 2], c=y, cmap='viridis', edgecolors='k')
ax.set_xlabel("Feature 1")
ax.set_ylabel("Feature 2")
ax.set_zlabel("Feature 3")
ax.set_title("Original Data (3D)")

# Step 3: Apply PCA (reduce to 2 components)
pca = PCA(n_components=2)
```

```
X_pca = pca.fit_transform(X)

# Step 4: Scatter Plot of Transformed 2D Data
ax2 = fig.add_subplot(122)
scatter2 = ax2.scatter(X_pca[:, 0], X_pca[:, 1], c=y, cmap='viridis', edgecolors='k')
ax2.set_xlabel("Principal Component 1")
ax2.set_ylabel("Principal Component 2")
ax2.set_title("Transformed Data after PCA (2D)")
plt.colorbar(scatter2, ax=ax2, label="Class Label")

plt.show()
```



```
In [26]: #Train a KNN Classifier on a real-world dataset (Wine dataset) and print classifica

from sklearn.datasets import load_wine
from sklearn.metrics import classification_report

# Step 1: Load the Wine dataset
wine = load_wine()
X = wine.data # Features
y = wine.target # Labels

# Step 2: Split dataset (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_sta

# Step 3: Train KNN Classifier (k=5)
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train, y_train)

# Step 4: Make Predictions
y_pred = knn.predict(X_test)

# Step 5: Print Classification Report
print("Classification Report:\n", classification_report(y_test, y_pred, target_name
```

Classification Report:

	precision	recall	f1-score	support
class_0	1.00	1.00	1.00	12
class_1	0.77	0.71	0.74	14
class_2	0.64	0.70	0.67	10
accuracy			0.81	36
macro avg	0.80	0.80	0.80	36
weighted avg	0.81	0.81	0.81	36

In [28]: *#Train a KNN Regressor and analyze the effect of different distance metrics on pred*

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.neighbors import KNeighborsRegressor
from sklearn.model_selection import train_test_split
from sklearn.datasets import load_wine
from sklearn.metrics import mean_squared_error

data = load_wine()
X, y = data.data, data.target

# Step 2: Split dataset (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Step 3: Train KNN Regressor with different distance metrics
distance_metrics = ['euclidean', 'manhattan', 'chebyshev', 'minkowski']
errors = []


for metric in distance_metrics:
    knn = KNeighborsRegressor(n_neighbors=5, metric=metric)
    knn.fit(X_train, y_train)
    y_pred = knn.predict(X_test)

    mse = mean_squared_error(y_test, y_pred)
    errors.append(mse)
    print(f"Metric: {metric} -> MSE: {mse:.4f}")

# Step 4: Plot Prediction Error vs. Distance Metric
plt.figure(figsize=(8, 5))
plt.bar(distance_metrics, errors, color=['b', 'g', 'r', 'c'])
plt.xlabel("Distance Metric")
plt.ylabel("Mean Squared Error (MSE)")
plt.title("Effect of Distance Metric on KNN Regression Error")
plt.show()

# Step 5: Print best metric
best_metric = distance_metrics[np.argmin(errors)]
print(f"\nBest distance metric: {best_metric} (Lowest MSE: {min(errors):.4f})")
```

```
Metric: euclidean -> MSE: 0.2900
Metric: manhattan -> MSE: 0.2167
Metric: chebyshev -> MSE: 0.2878
Metric: minkowski -> MSE: 0.2900
```



Best distance metric: manhattan (Lowest MSE: 0.2167)

```
In [32]: #Train a KNN Classifier and evaluate using ROC-AUC score
import numpy as np
import matplotlib.pyplot as plt
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
from sklearn.datasets import load_wine
from sklearn.metrics import roc_auc_score, roc_curve
from sklearn.preprocessing import label_binarize

# Load dataset
wine = load_wine()
X = wine.data
y = wine.target

# Convert to binary format (One-vs-Rest)
y_bin = label_binarize(y, classes=[0, 1, 2])

# Split dataset
X_train, X_test, y_train, y_test = train_test_split(X, y_bin, test_size=0.2, random

# Train KNN Classifier
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train, y_train)

# Get predicted probabilities
y_scores = np.array(knn.predict_proba(X_test)) # Convert to NumPy array

# Debugging: Print the shape
print("y_scores shape:", y_scores.shape) # Should match y_test shape
```



```

# Compute AUC for each class
auc_scores = []
plt.figure(figsize=(8, 6))

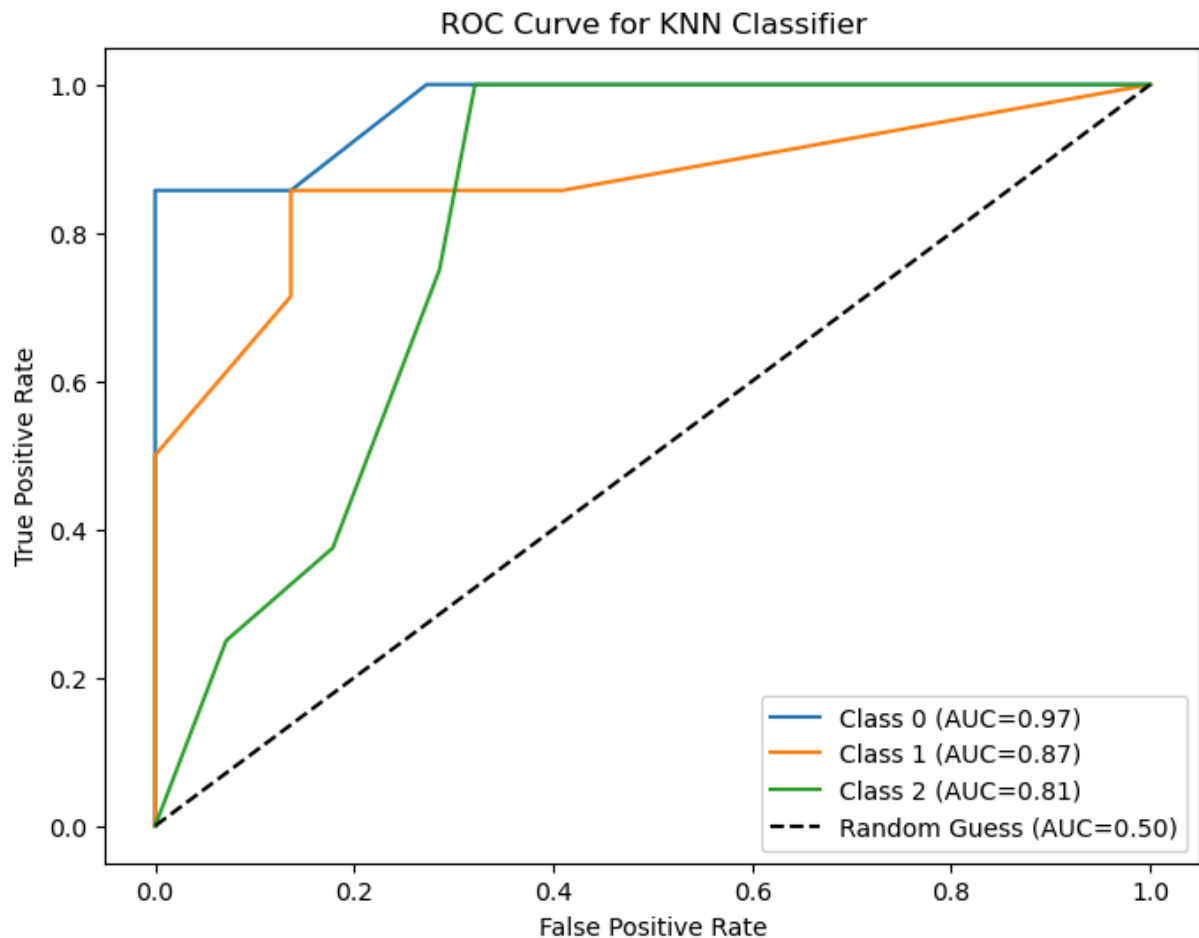
for i in range(y_bin.shape[1]):
    fpr, tpr, _ = roc_curve(y_test[:, i], y_scores[i][:, 1]) # Fix indexing
    auc = roc_auc_score(y_test[:, i], y_scores[i][:, 1])
    auc_scores.append(auc)
    plt.plot(fpr, tpr, label=f'Class {i} (AUC={auc:.2f})')

# Plot ROC Curve
plt.plot([0, 1], [0, 1], 'k--', label="Random Guess (AUC=0.50)")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve for KNN Classifier")
plt.legend()
plt.show()

# Print AUC scores
print(f"Mean AUC Score: {np.mean(auc_scores):.4f}")

```

y_scores shape: (3, 36, 2)



Mean AUC Score: 0.8828

```

In [33]: #Train a PCA model and visualize the variance captured by each principal component
import numpy as np
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA

```

```

from sklearn.preprocessing import StandardScaler
from sklearn.datasets import load_wine

# Load dataset
wine = load_wine()
X = wine.data

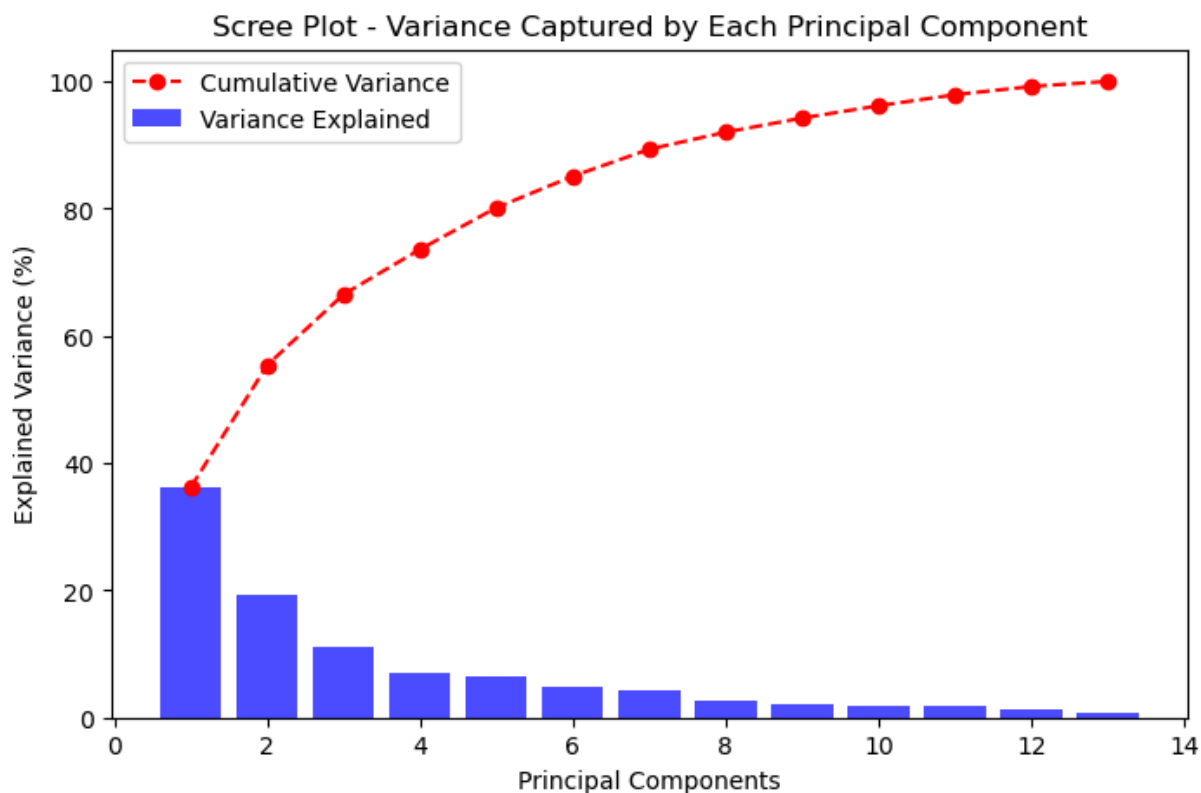
# Standardize the data (PCA is sensitive to scale)
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Train PCA model
pca = PCA(n_components=X.shape[1]) # Keep all components
X_pca = pca.fit_transform(X_scaled)

# Get explained variance ratio
explained_variance = pca.explained_variance_ratio_

# Plot Scree plot
plt.figure(figsize=(8, 5))
plt.bar(range(1, len(explained_variance) + 1), explained_variance * 100, alpha=0.7,
plt.plot(range(1, len(explained_variance) + 1), np.cumsum(explained_variance) * 100
plt.xlabel("Principal Components")
plt.ylabel("Explained Variance (%)")
plt.title("Scree Plot - Variance Captured by Each Principal Component")
plt.legend()
plt.show()

```



```

In [34]: #Train a KNN Classifier and perform feature selection before training
import numpy as np
import matplotlib.pyplot as plt

```

```

from sklearn.datasets import load_wine
from sklearn.preprocessing import StandardScaler
from sklearn.feature_selection import SelectKBest, f_classif
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score

# Load dataset
wine = load_wine()
X, y = wine.data, wine.target

# Standardize the data (KNN is distance-based, so scaling is necessary)
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Feature Selection - Select Top K Features
k = 5 # Choose top 5 best features
selector = SelectKBest(score_func=f_classif, k=k)
X_selected = selector.fit_transform(X_scaled, y)

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_selected, y, test_size=0.2, r

# Train KNN Classifier
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train, y_train)

# Predict & Evaluate
y_pred = knn.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)

# Print Results
print(f"Selected Features: {np.array(wine.feature_names)[selector.get_support()]}")
print(f"Accuracy with Feature Selection: {accuracy:.4f}")

```

Selected Features: ['alcohol' 'flavanoids' 'color_intensity' 'od280/od315_of_diluted_wines' 'proline']
 Accuracy with Feature Selection: 0.9722

In [35]: *#Train a PCA model and visualize the data reconstruction error after reducing dimen*

```

import numpy as np
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.datasets import load_wine
from sklearn.metrics import mean_squared_error

# Load dataset
wine = load_wine()
X = wine.data


# Standardize features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

```

```
# Try different numbers of PCA components
components_range = range(1, X.shape[1] + 1)
reconstruction_errors = []

for n_components in components_range:
    pca = PCA(n_components=n_components)
    X_pca = pca.fit_transform(X_scaled) # Reduce dimensions
    X_reconstructed = pca.inverse_transform(X_pca) # Reconstruct data
    error = mean_squared_error(X_scaled, X_reconstructed) # Compute reconstruction error
    reconstruction_errors.append(error)

# Plot Reconstruction Error vs. Number of Components
plt.figure(figsize=(8, 5))
plt.plot(components_range, reconstruction_errors, marker='o', linestyle='--', color='blue')
plt.xlabel('Number of Principal Components')
plt.ylabel('Reconstruction Error (MSE)')
plt.title('PCA Data Reconstruction Error')
plt.grid()
plt.show()
```



```
In [36]: #Train a KNN Classifier and visualize the decision boundary
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import make_moons
from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import StandardScaler

# Generate synthetic dataset
X, y = make_moons(n_samples=300, noise=0.2, random_state=42)

# Standardize features
```

```

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)


# Train KNN Classifier
knn = KNeighborsClassifier(n_neighbors=5, weights='uniform')
knn.fit(X_scaled, y)

# Create mesh grid for visualization
x_min, x_max = X_scaled[:, 0].min() - 0.5, X_scaled[:, 0].max() + 0.5
y_min, y_max = X_scaled[:, 1].min() - 0.5, X_scaled[:, 1].max() + 0.5
xx, yy = np.meshgrid(np.linspace(x_min, x_max, 100), np.linspace(y_min, y_max, 100))

# Predict on grid points
Z = knn.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

# Plot decision boundary
plt.figure(figsize=(8, 6))
plt.contourf(xx, yy, Z, alpha=0.3, cmap=plt.cm.coolwarm)
plt.scatter(X_scaled[:, 0], X_scaled[:, 1], c=y, edgecolors='k', cmap=plt.cm.coolwa)
plt.xlabel("Feature 1")
plt.ylabel("Feature 2")
plt.title("KNN Decision Boundary (k=5)")
plt.show()

```



In [37]: # Train a PCA model and analyze the effect of different numbers of components on da
import numpy as np

```
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA
from sklearn.datasets import load_digits
from sklearn.preprocessing import StandardScaler


# Load dataset (digits dataset with 64 features)
digits = load_digits()
X = digits.data

# Standardize the dataset
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Apply PCA with varying number of components
n_components = np.arange(1, X.shape[1] + 1)
explained_variances = []

for n in n_components:
    pca = PCA(n_components=n)
    pca.fit(X_scaled)
    explained_variances.append(np.sum(pca.explained_variance_ratio_)) # Cumulative

# Plot explained variance vs. number of components
plt.figure(figsize=(8, 6))
plt.plot(n_components, explained_variances, marker='o', linestyle='-', color='b')
plt.xlabel("Number of Principal Components")
plt.ylabel("Cumulative Explained Variance")
plt.title("Effect of PCA Components on Data Variance")
plt.grid(True)
plt.axhline(y=0.95, color='r', linestyle='--', label='95% Variance')
plt.legend()
plt.show()
```



In []: